

Building event-driven (Micro)Services with Apache Kafka

Guido Schmutz

Voxxed Days Banff – 27.10.2018

VOXXED DAYS

BANFF

26th - 27th October, 2018



@gschmutz



guidoschmutz.wordpress.com

BASEL ▪ BERN ▪ BRUGG ▪ DÜSSELDORF ▪ FRANKFURT A.M. ▪ FREIBURG I.BR. ▪ GENF
HAMBURG ▪ KOPENHAGEN ▪ LAUSANNE ▪ MÜNCHEN ▪ STUTTGART ▪ WIEN ▪ ZÜRICH

trivadis
makes **IT** easier. 

■ Guido Schmutz

Working at Trivadis for more than 21 years

Oracle ACE Director for Fusion Middleware and SOA

Consultant, Trainer Software Architect for Java, Oracle, SOA and
Big Data / Fast Data

Head of Trivadis Architecture Board

Technology Manager @ Trivadis

More than 30 years of software development experience

Contact: guido.schmutz@trivadis.com

Blog: <http://guidoschmutz.wordpress.com>

Slideshare: <http://www.slideshare.net/gschmutz>

Twitter: [@gschmutz](https://twitter.com/gschmutz)



gechmutz 24.08.17 on April 18, 2017

Tags: flink (1), kafka (59), kafka-connect (4), kafka-streams (17), spark-streaming (31), storm (39), streams (4)

Last week in Stream Processing & Analytics – 18.4.2017

This is the 62nd edition of my blog series blog series around Stream Processing and Analytics!

Every week I'm also updating the following two lists with the presentations/videos of the current week:

- [Presentations from Slideshare](#)
- [Videos from YouTube](#)

As usual, find below the new blog articles, presentations, videos and software releases from last week:

News and Blog Posts

General

- [Multi Master Replication For Geo-Distributed Data: It's more than you think](#) by Ellen Friedman
- [Understanding Indicators of Attack \(IOAs\): The Power of Event Stream Processing in CrowdStrike Falcon](#) by Dan Brown
- [Stream processing and messaging systems for the IoT age](#) by Ben Lorica

Apache Kafka / Kafka Streams / Confluent Platform

- [Creating a Data Pipeline with Kafka Connect API - from Architecture to Operations](#) by Alexandra Wang
- [Streaming Spring Boot Application Logs to ELK Stack—Part 1](#) by kaadayamuthu
- [Streaming Spring Boot Application Logs to Apache Kafka—ELK/Kafka Stack—Part 2](#) by kaadayamuthu



With over 600 specialists and IT experts in your region.



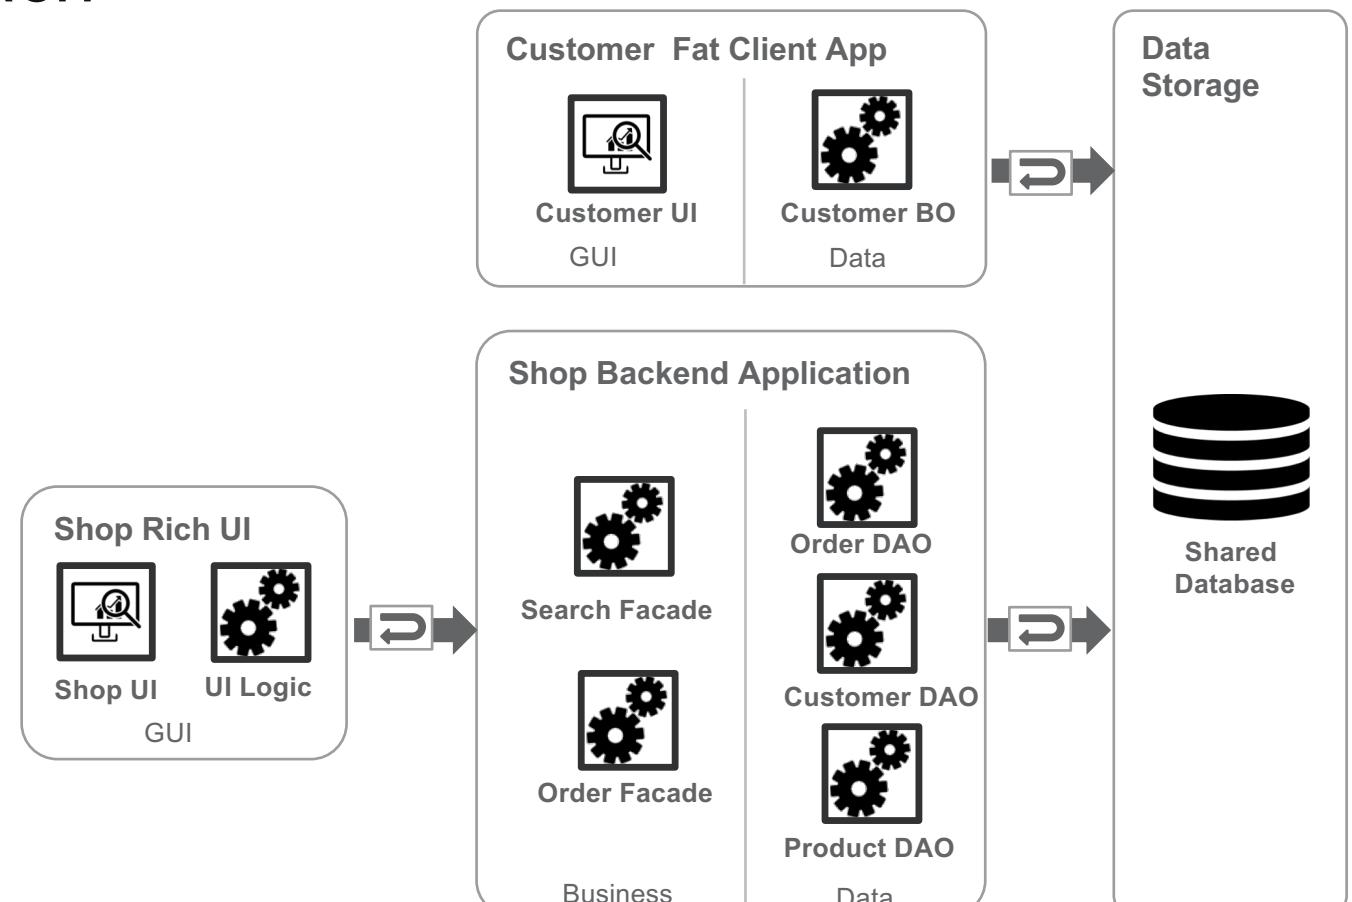
- 14 Trivadis branches and more than 600 employees
- 200 Service Level Agreements
- Over 4,000 training participants
- Research and development budget: CHF 5.0 million
- Financially self-supporting and sustainably profitable
- Experience from more than 1,900 projects per year at over 800 customers

■ Agenda

1. Where do we come from?
2. Using Microservices
3. Can we do (even) better?
4. Kafka for Event-Driven (Micro)Services
5. What about streaming sources?
6. What about integrating legacy applications?
7. What about (historical) data analytics?
8. Summary

Where do we come from?

■ Traditional Approach



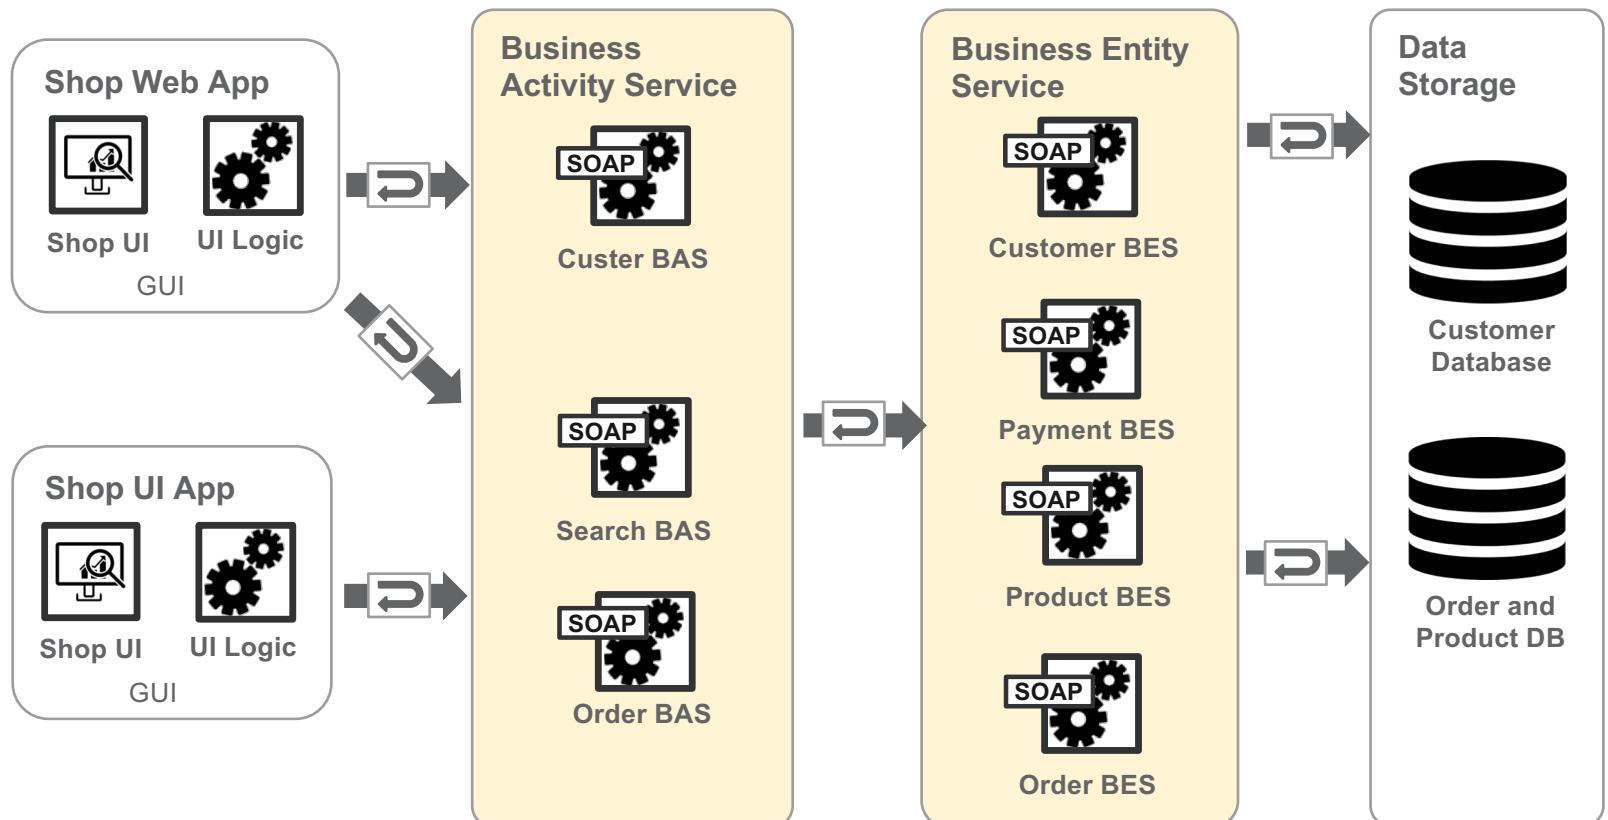
■ SOA Approach

Contract-first
Web Services

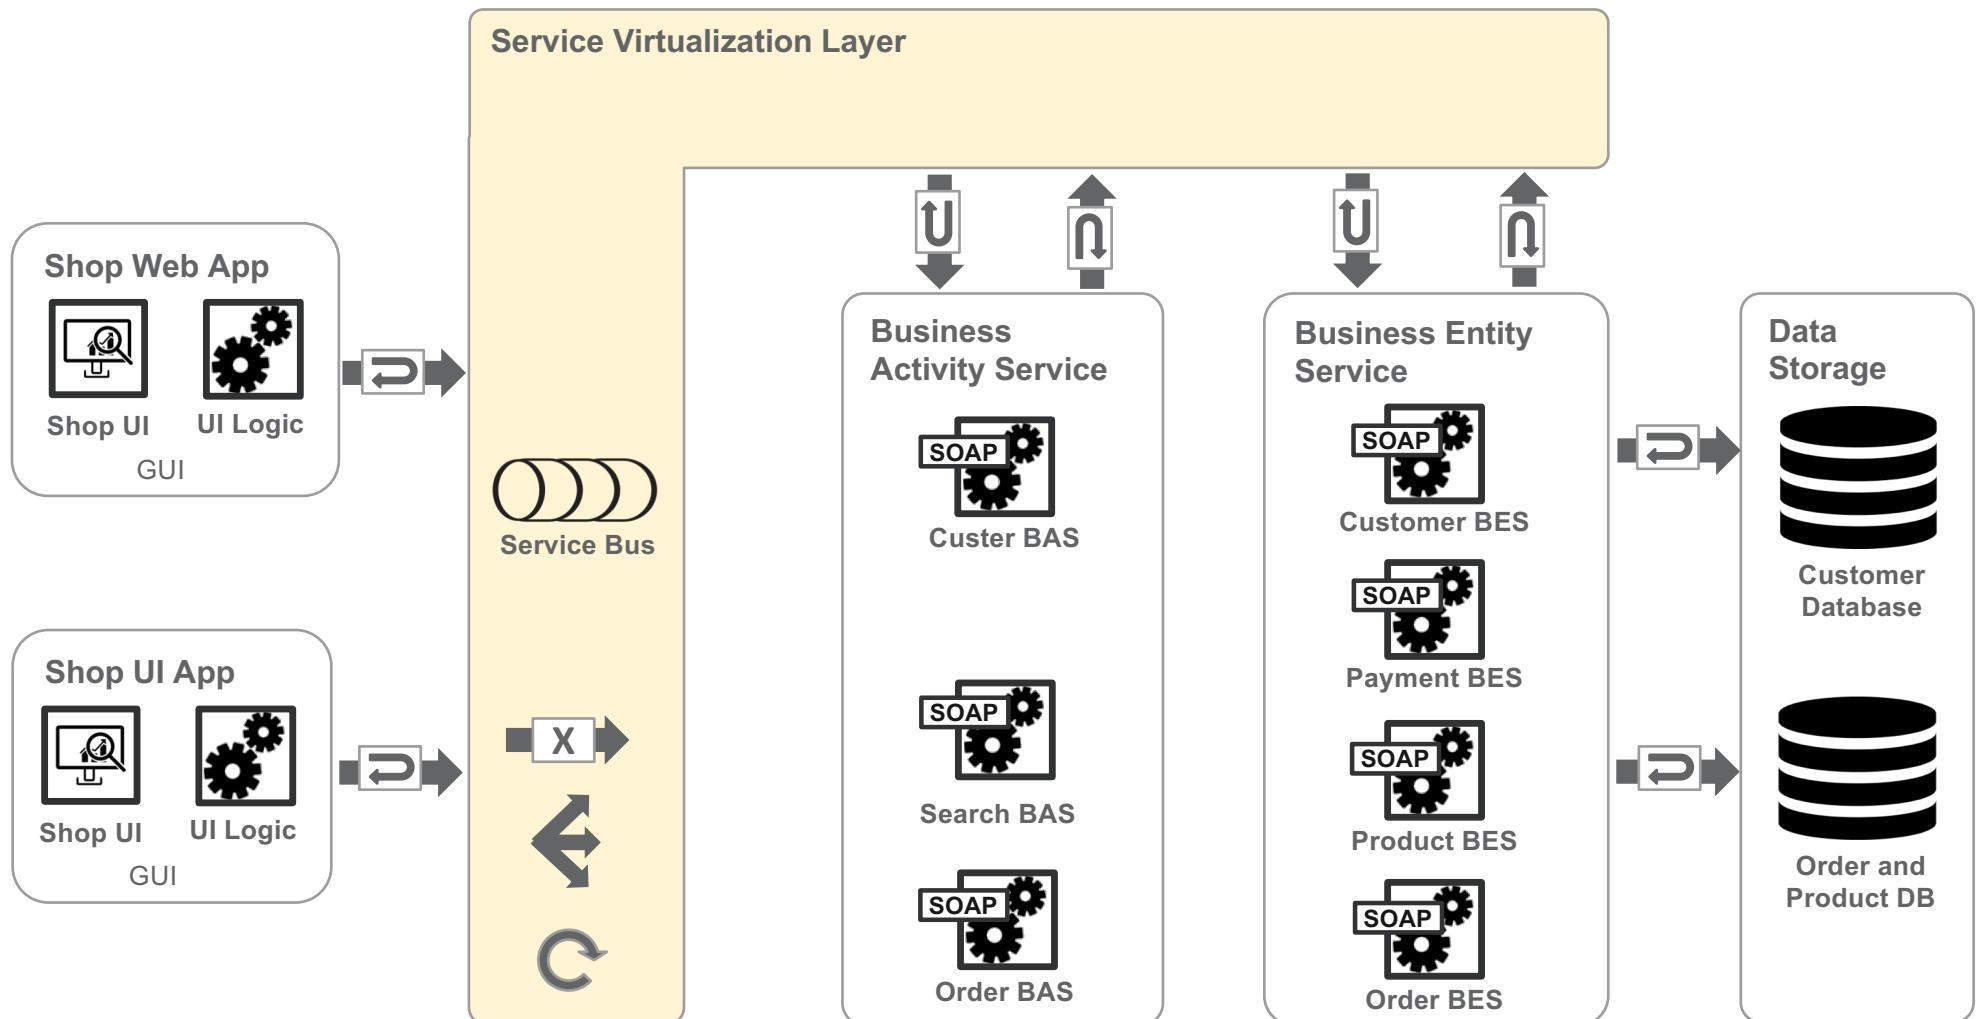
Technical layers
offer their own
interfaces

Reuse on each
level

Lower layer
often wraps
legacy code



Virtualized SOA Approach



Using Microservices

■ Microservice Approach

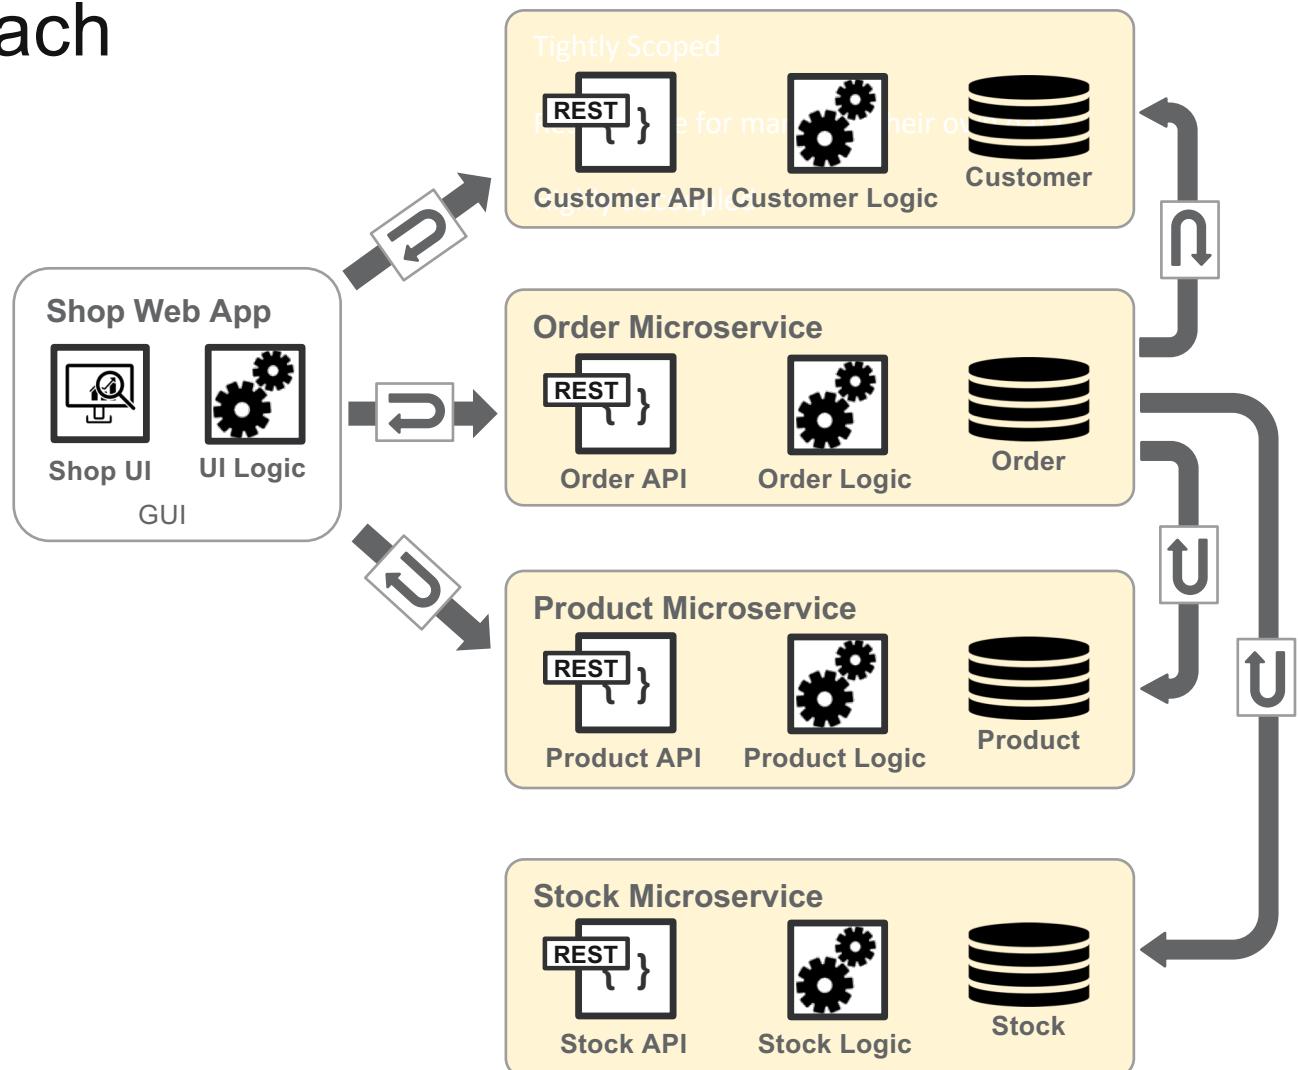
Tightly Scoped behind clear interfaces

Responsible for managing their own data (not necessarily the infrastructure)

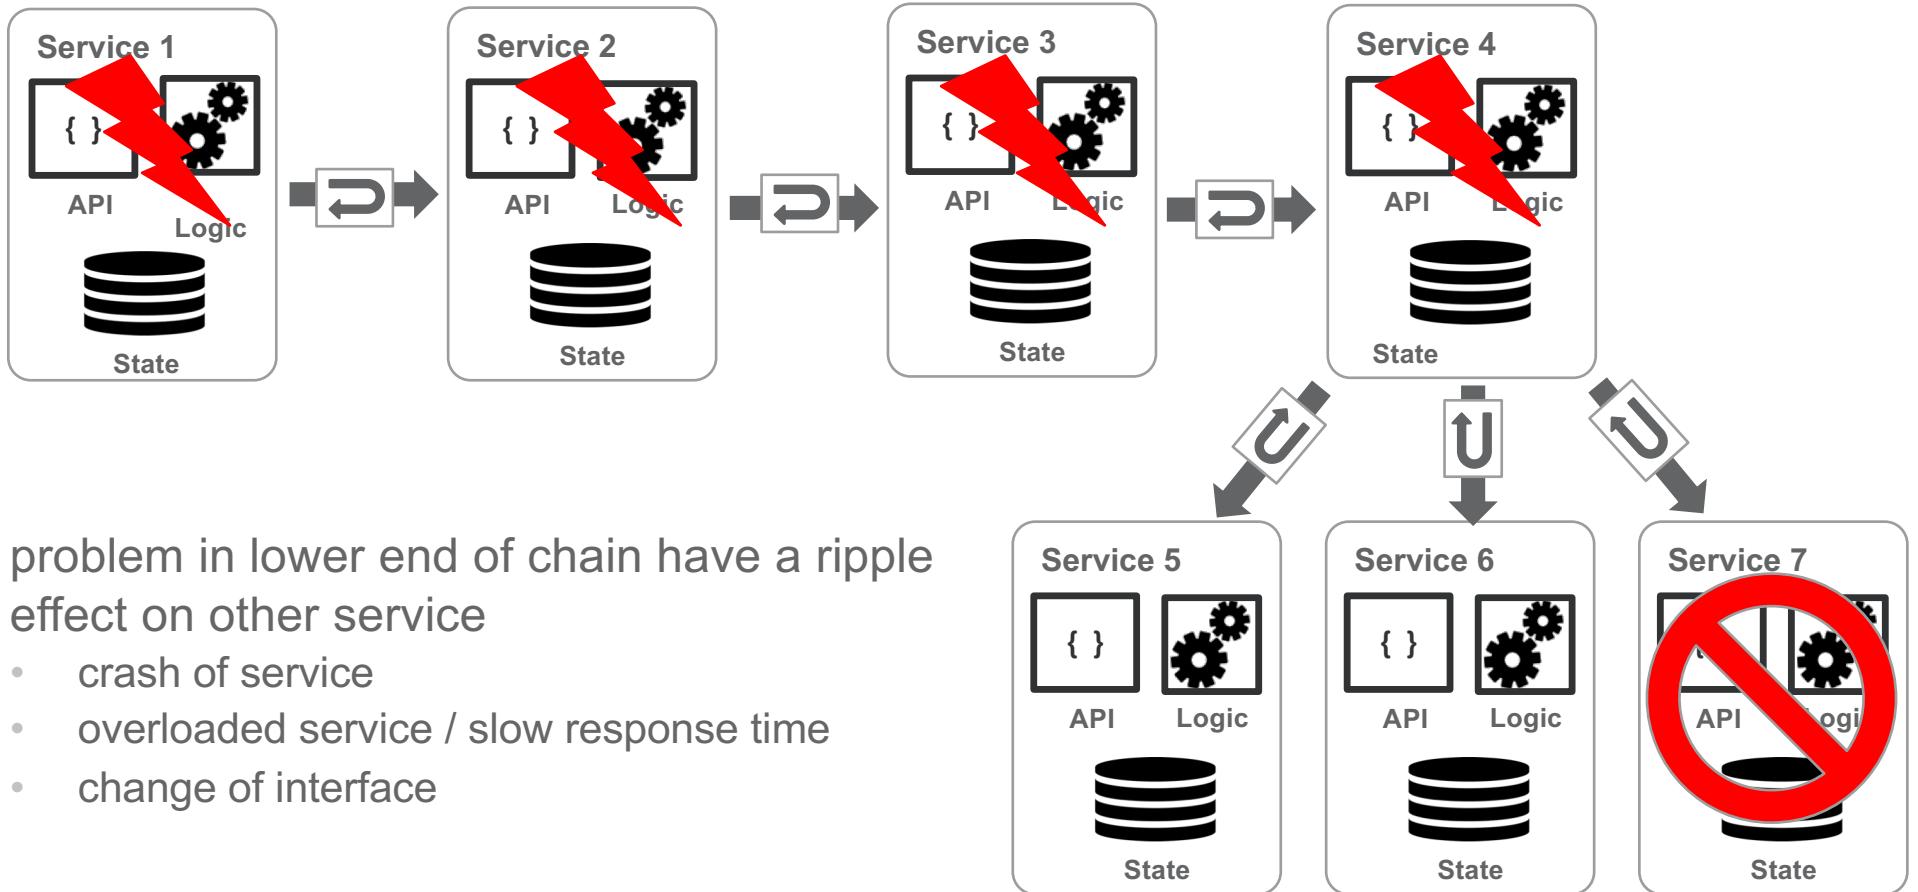
Should be highly decoupled

Independently deployable, self-contained and autonomous

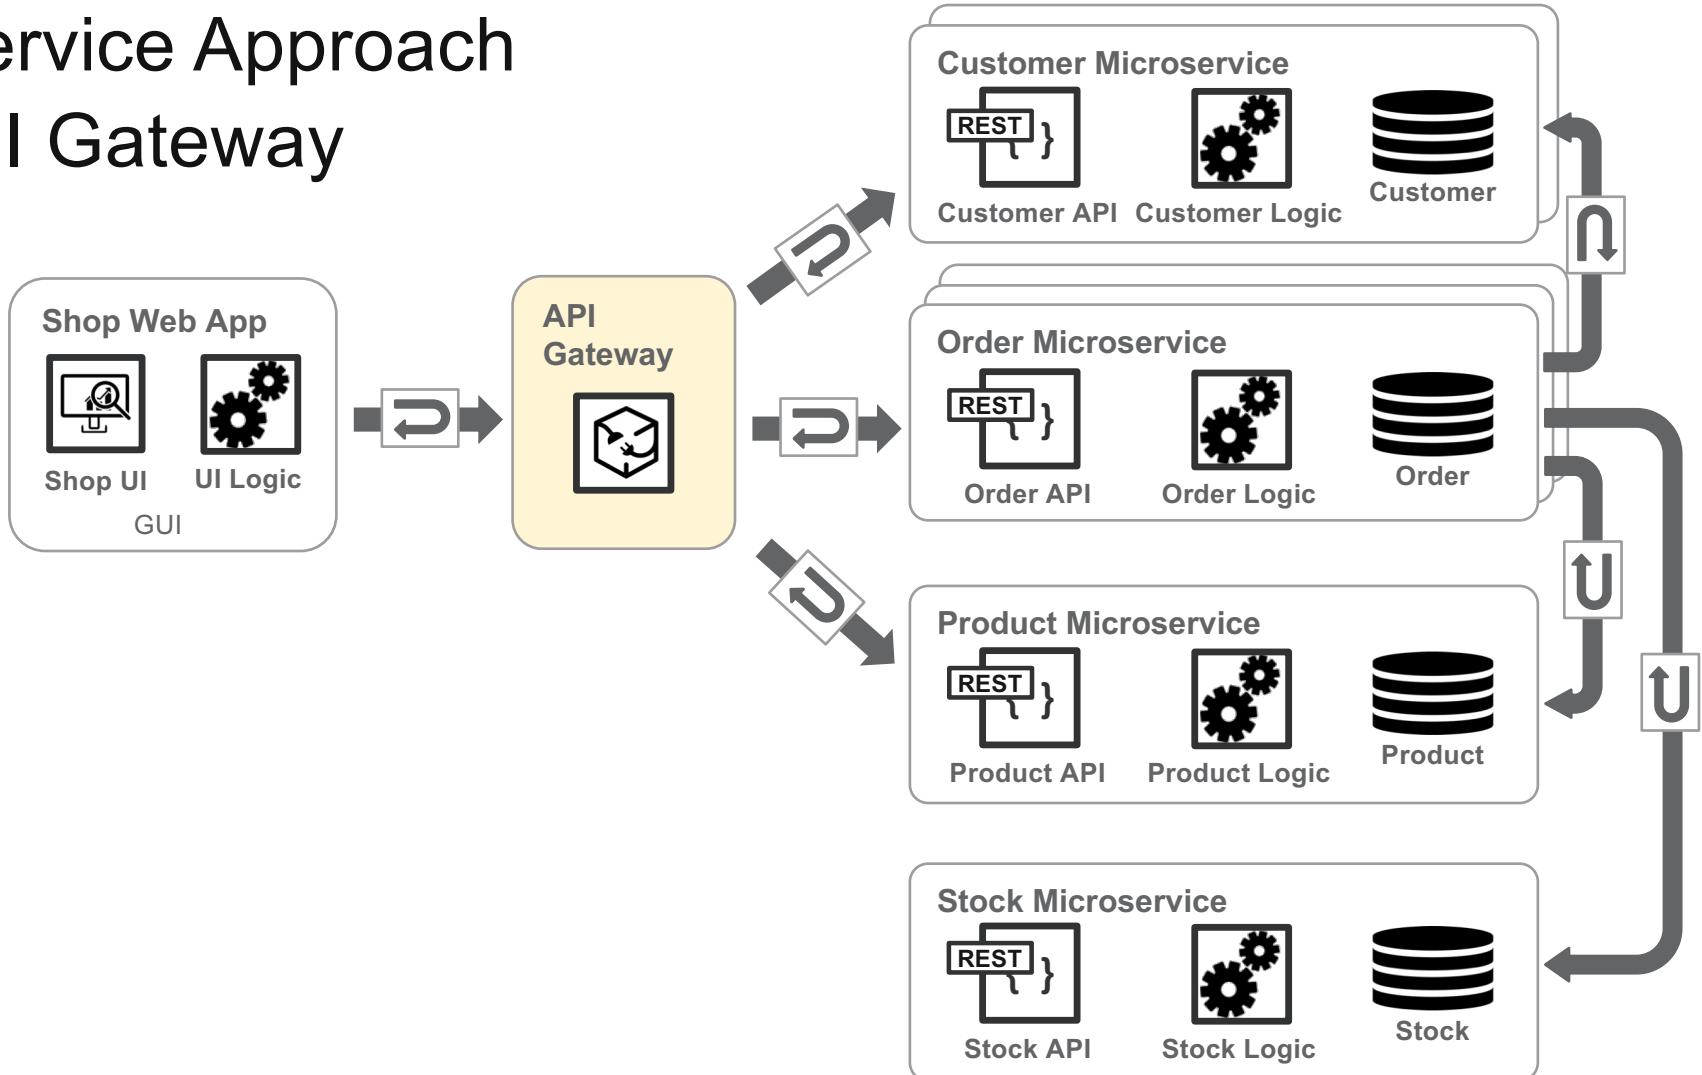
SOA done right ?!



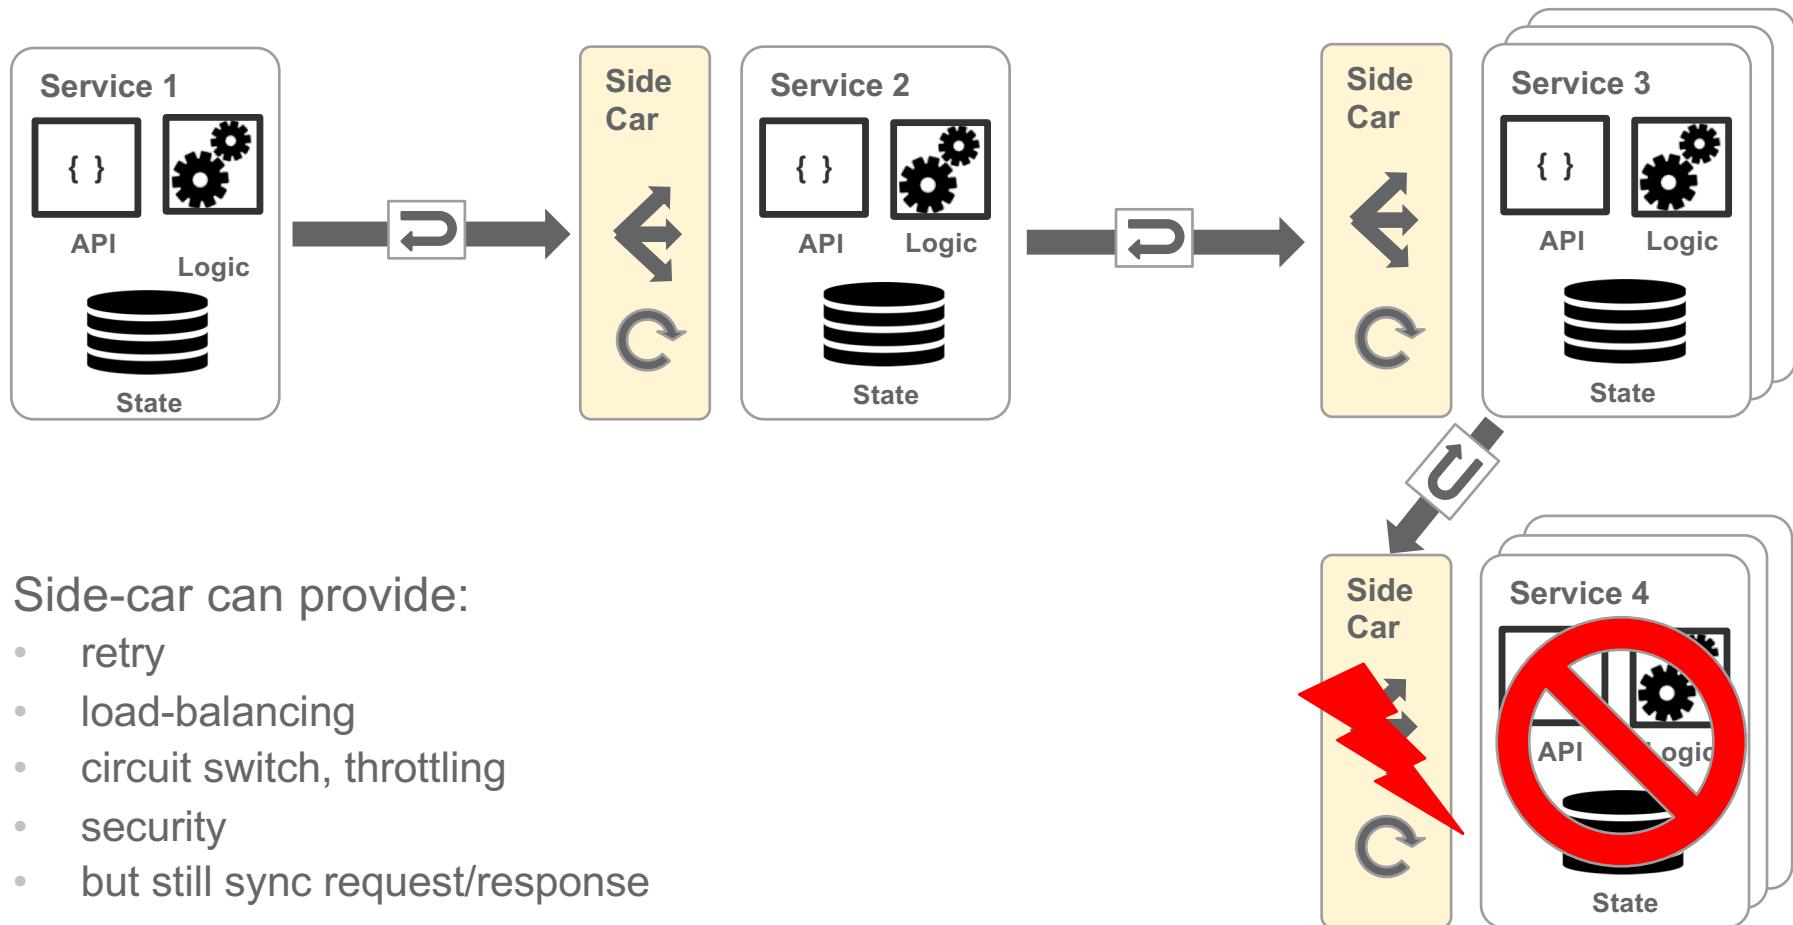
■ Synchronous World of Request-Response leads to tight, point-to-point couplings



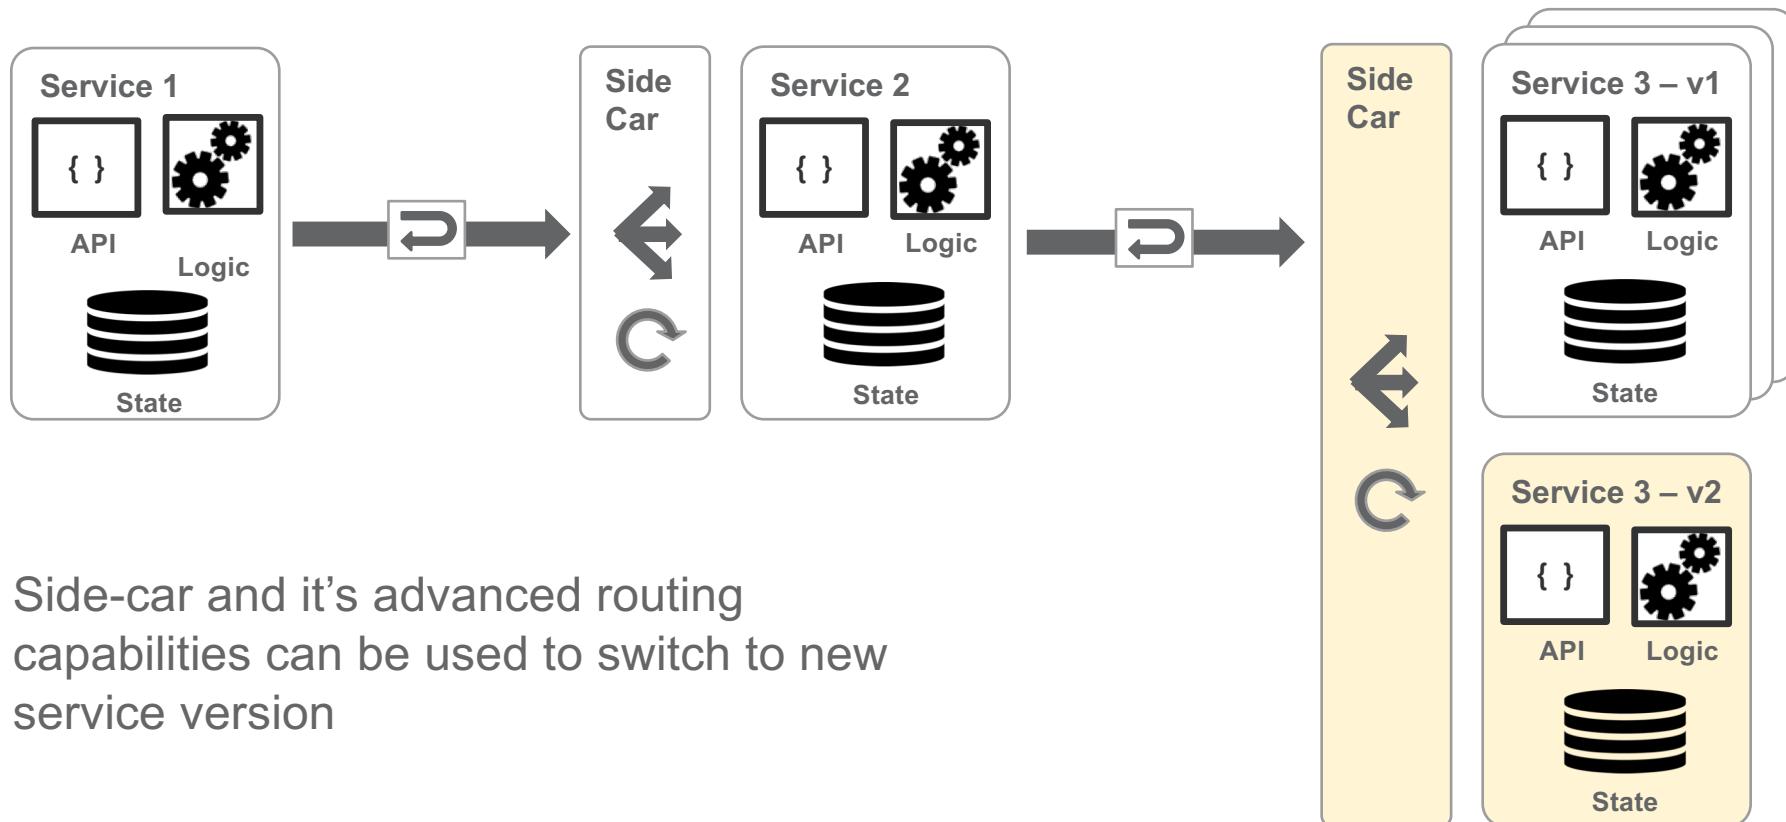
■ Microservice Approach with API Gateway



■ Microservice Approach with Side Car (i.e. K8s & Istio)

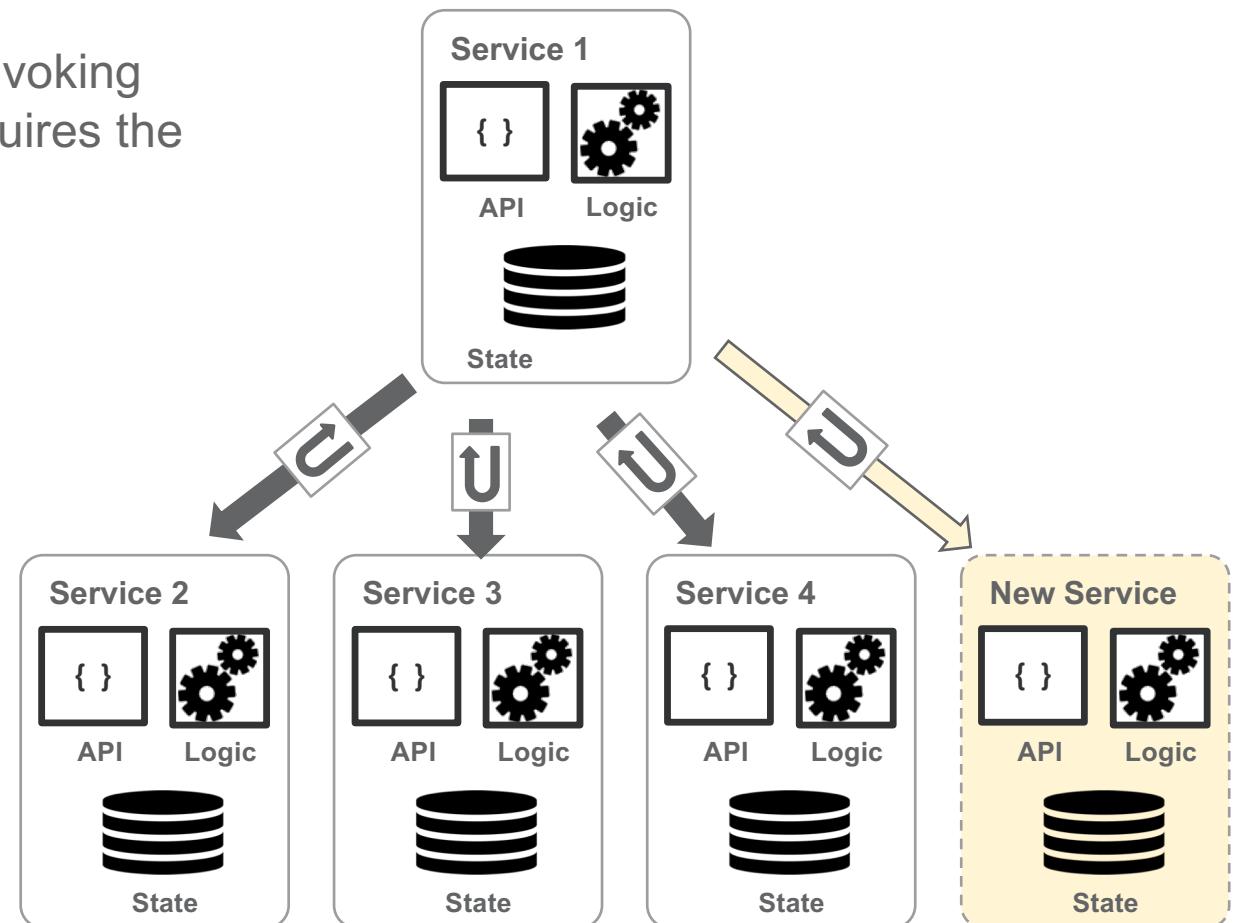


■ Microservice Approach with Side Car (i.e. K8s & Istio)



■ Side Car provides abstraction

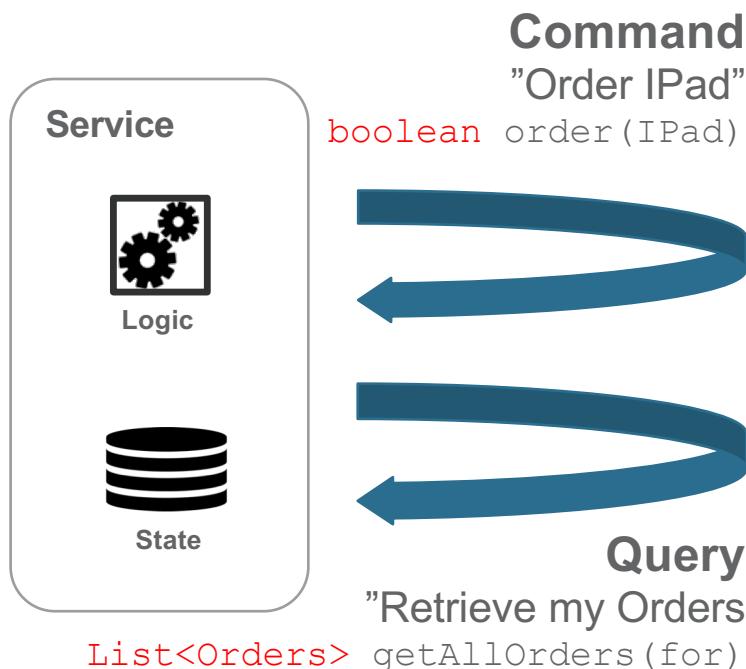
but we still have to change the invoking service If another service requires the same information



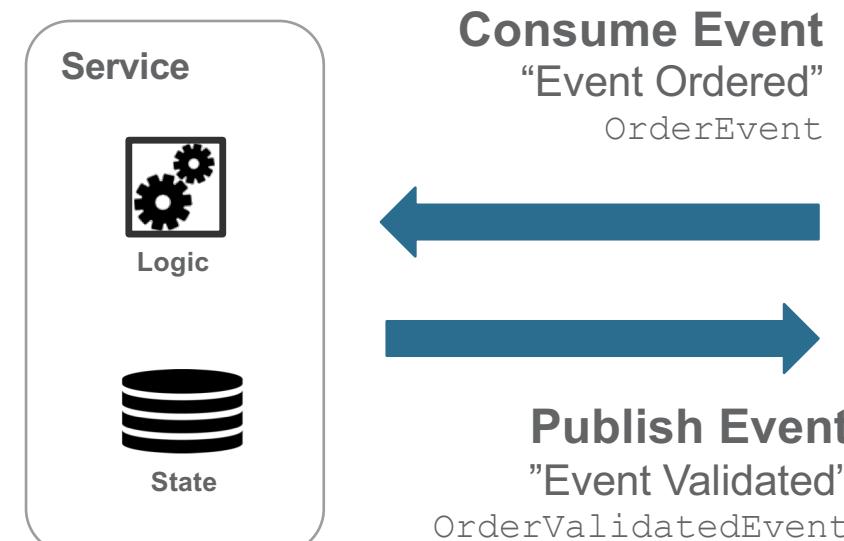
Can we do (even) better?

■ Three mechanisms through which services can interact

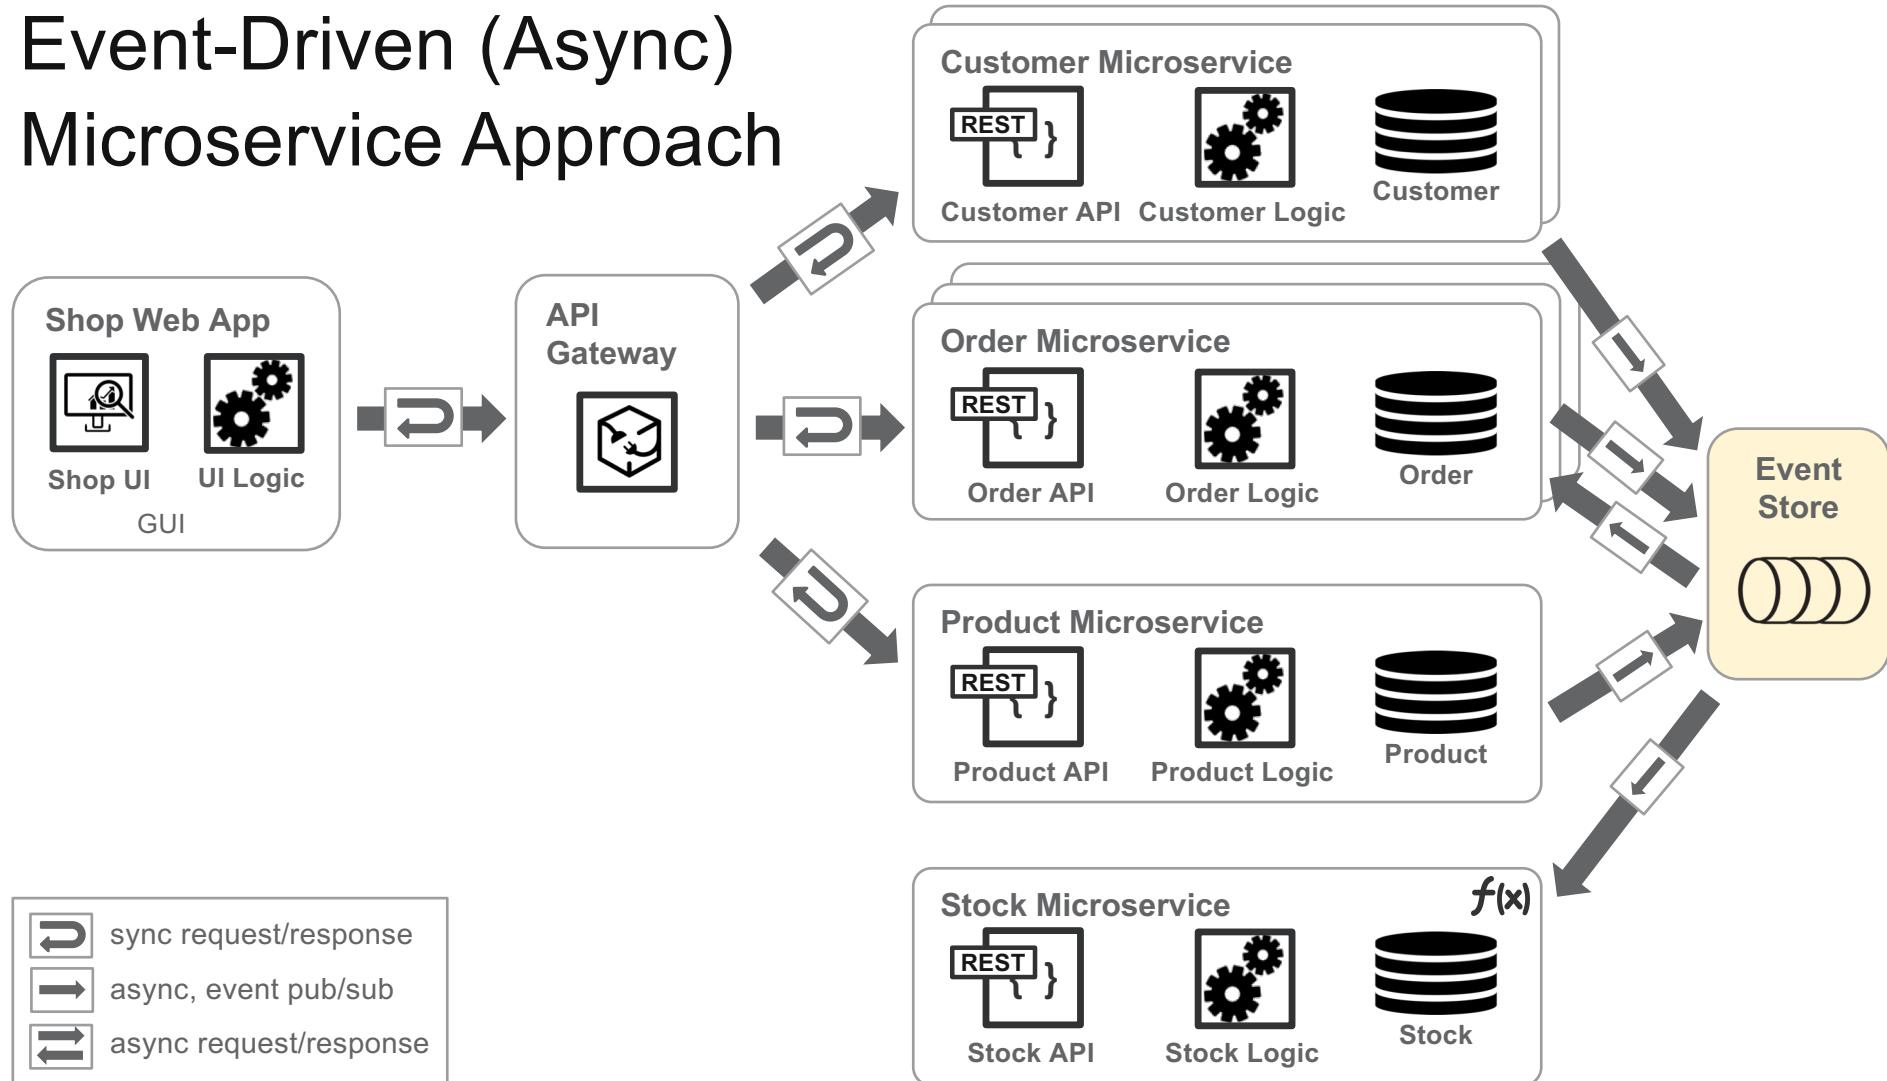
Request-Driven (Imperative)



Event Driven (Functional)



■ Event-Driven (Async) Microservice Approach

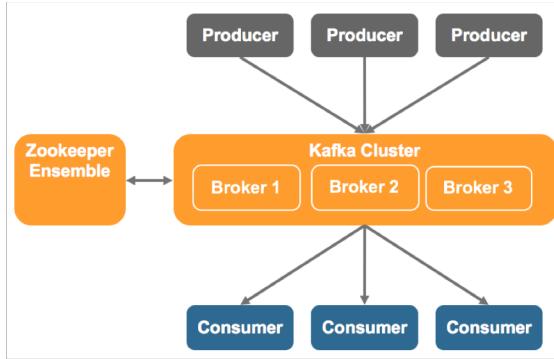


Kafka for Event-Driven (Micro)Services

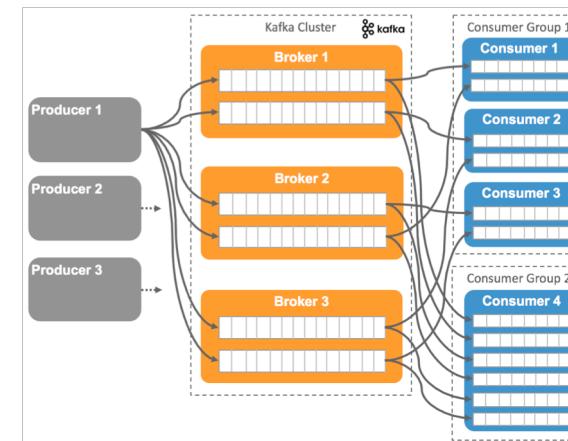
Apache Kafka – A Streaming Platform



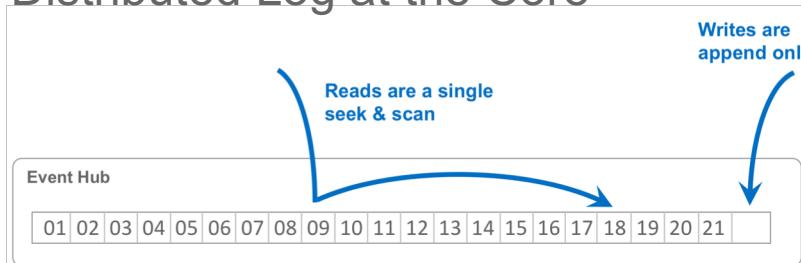
High-Level Architecture



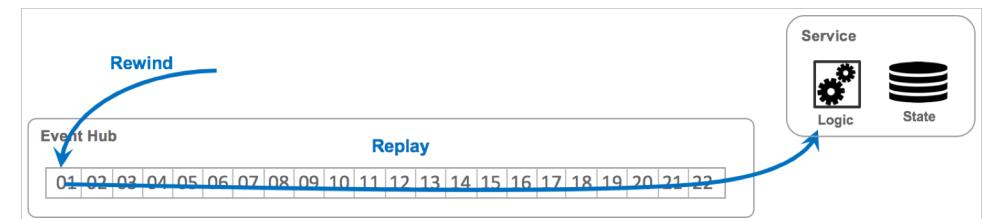
Scale-Out Architecture



Distributed Log at the Core



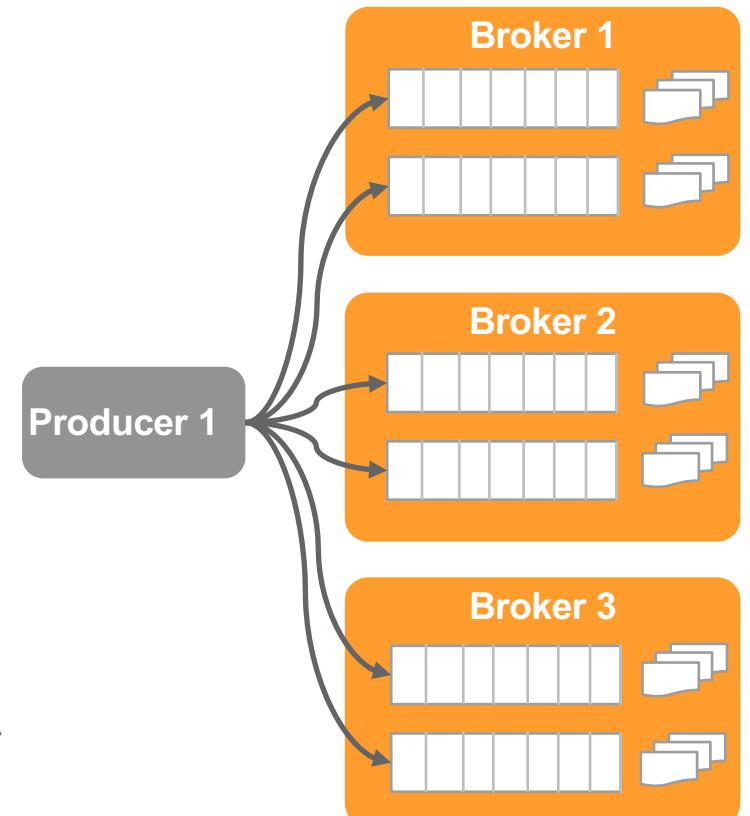
Logs do not (necessarily) forget



■ Hold Data for Long-Term – Data Retention

1. Never
2. Time based (TTL)
`log.retention.{ms | minutes | hours}`
3. Size based
`log.retention.bytes`
4. Log compaction based
(entries with same key are removed):

```
kafka-topics.sh --zookeeper zk:2181 \
    --create --topic customers \
    --replication-factor 1 \
    --partitions 1 \
    --config cleanup.policy=compact
```

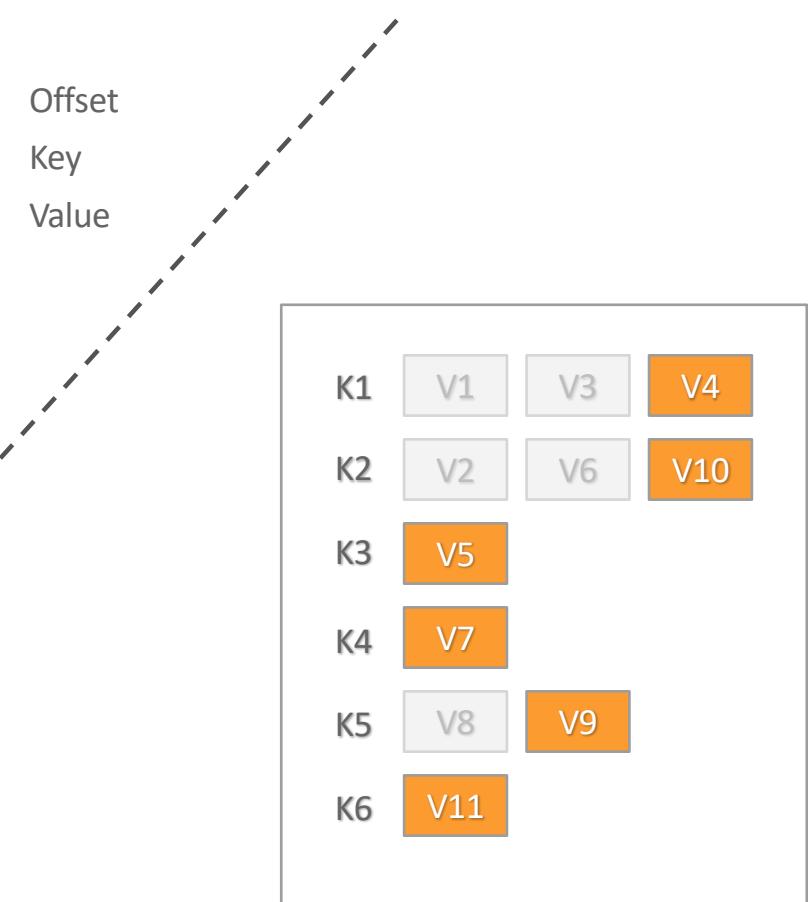


Keep Topics in Compacted Form

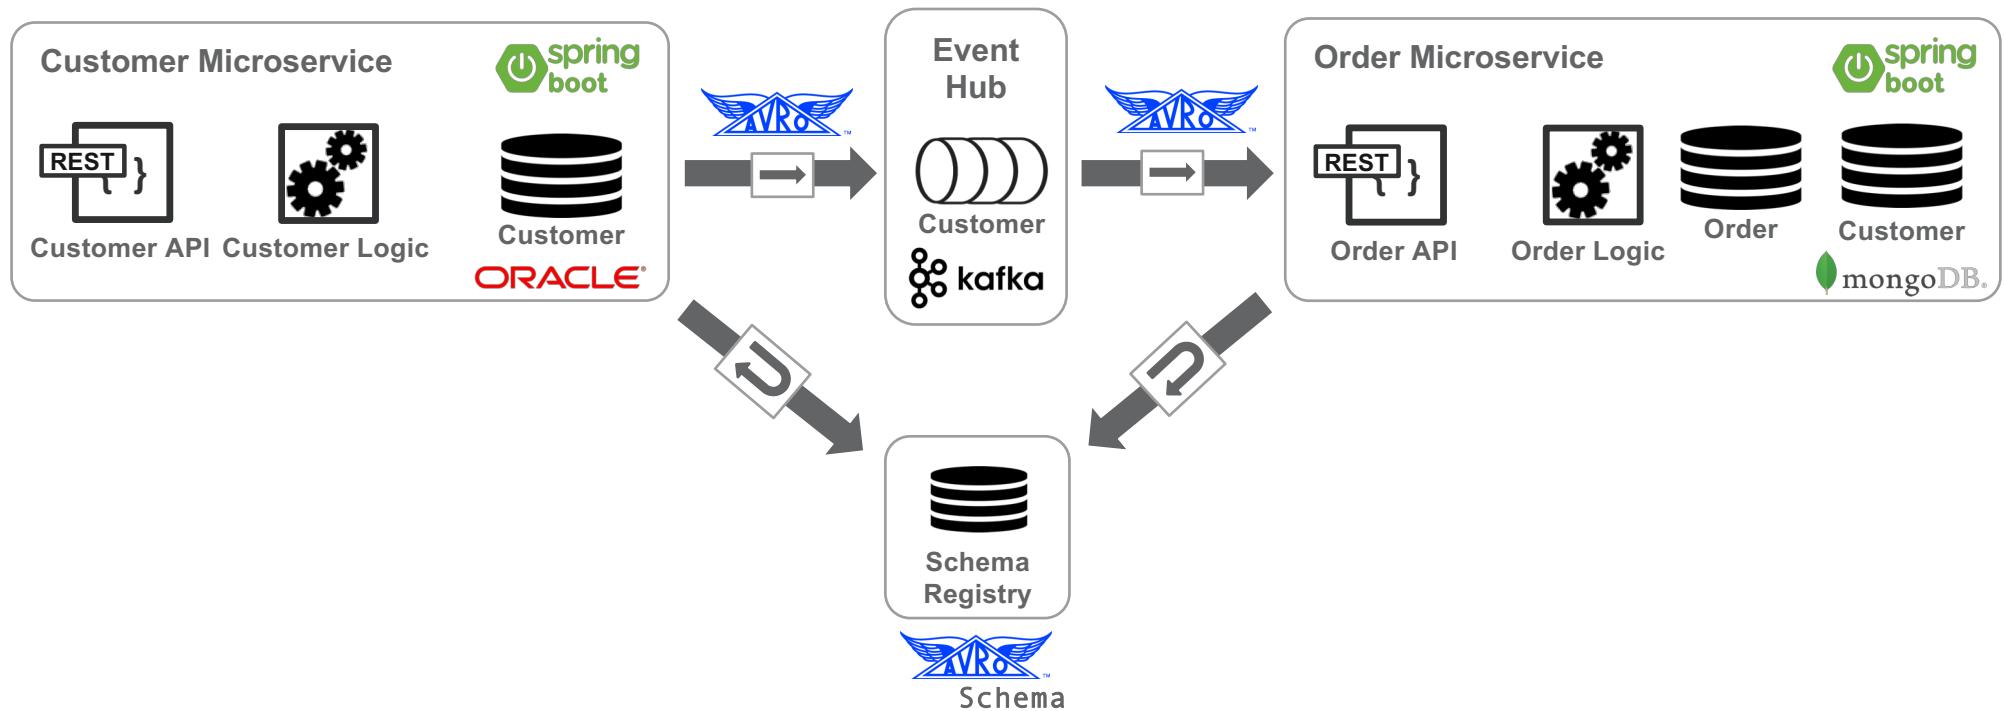
0	1	2	3	4	5	6	7	8	9	10	11
K1	K2	K1	K1	K3	K2	K4	K5	K5	K2	K6	K2
V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	

Compaction

Offset	3	4	6	8	9	10
Key	K1	K3	K4	K5	K2	K6
Value	V4	V5	V7	V9	V10	V11

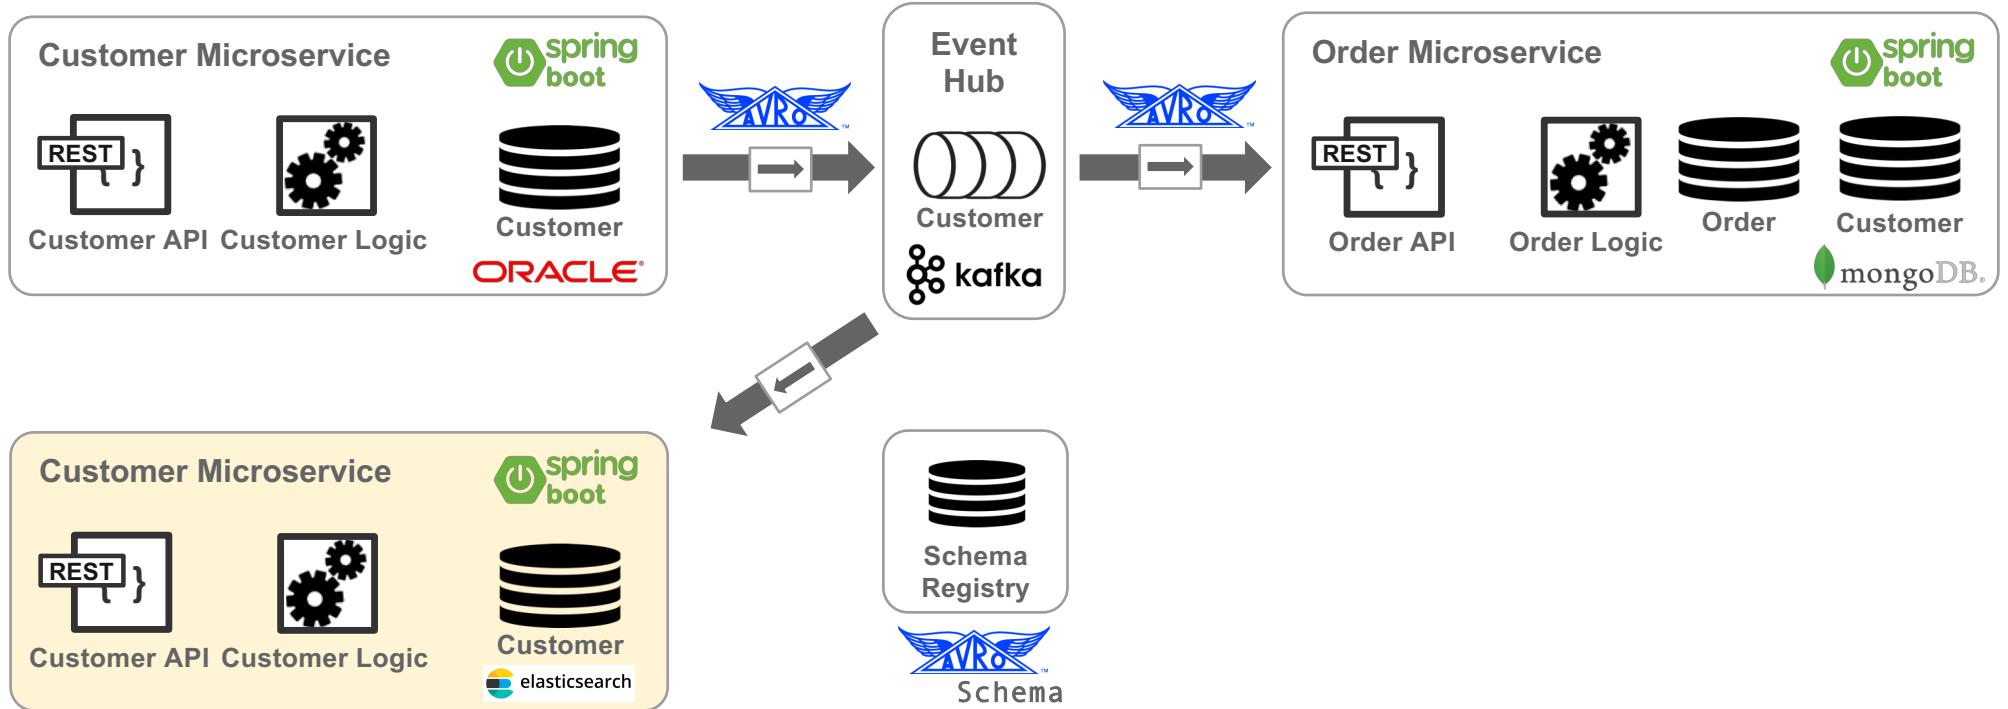
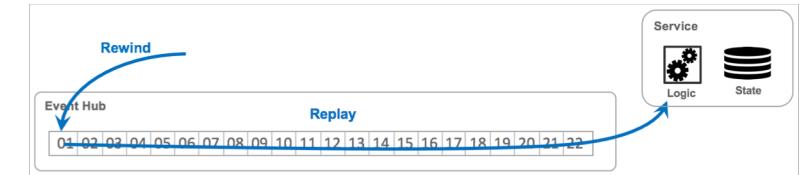


Demo



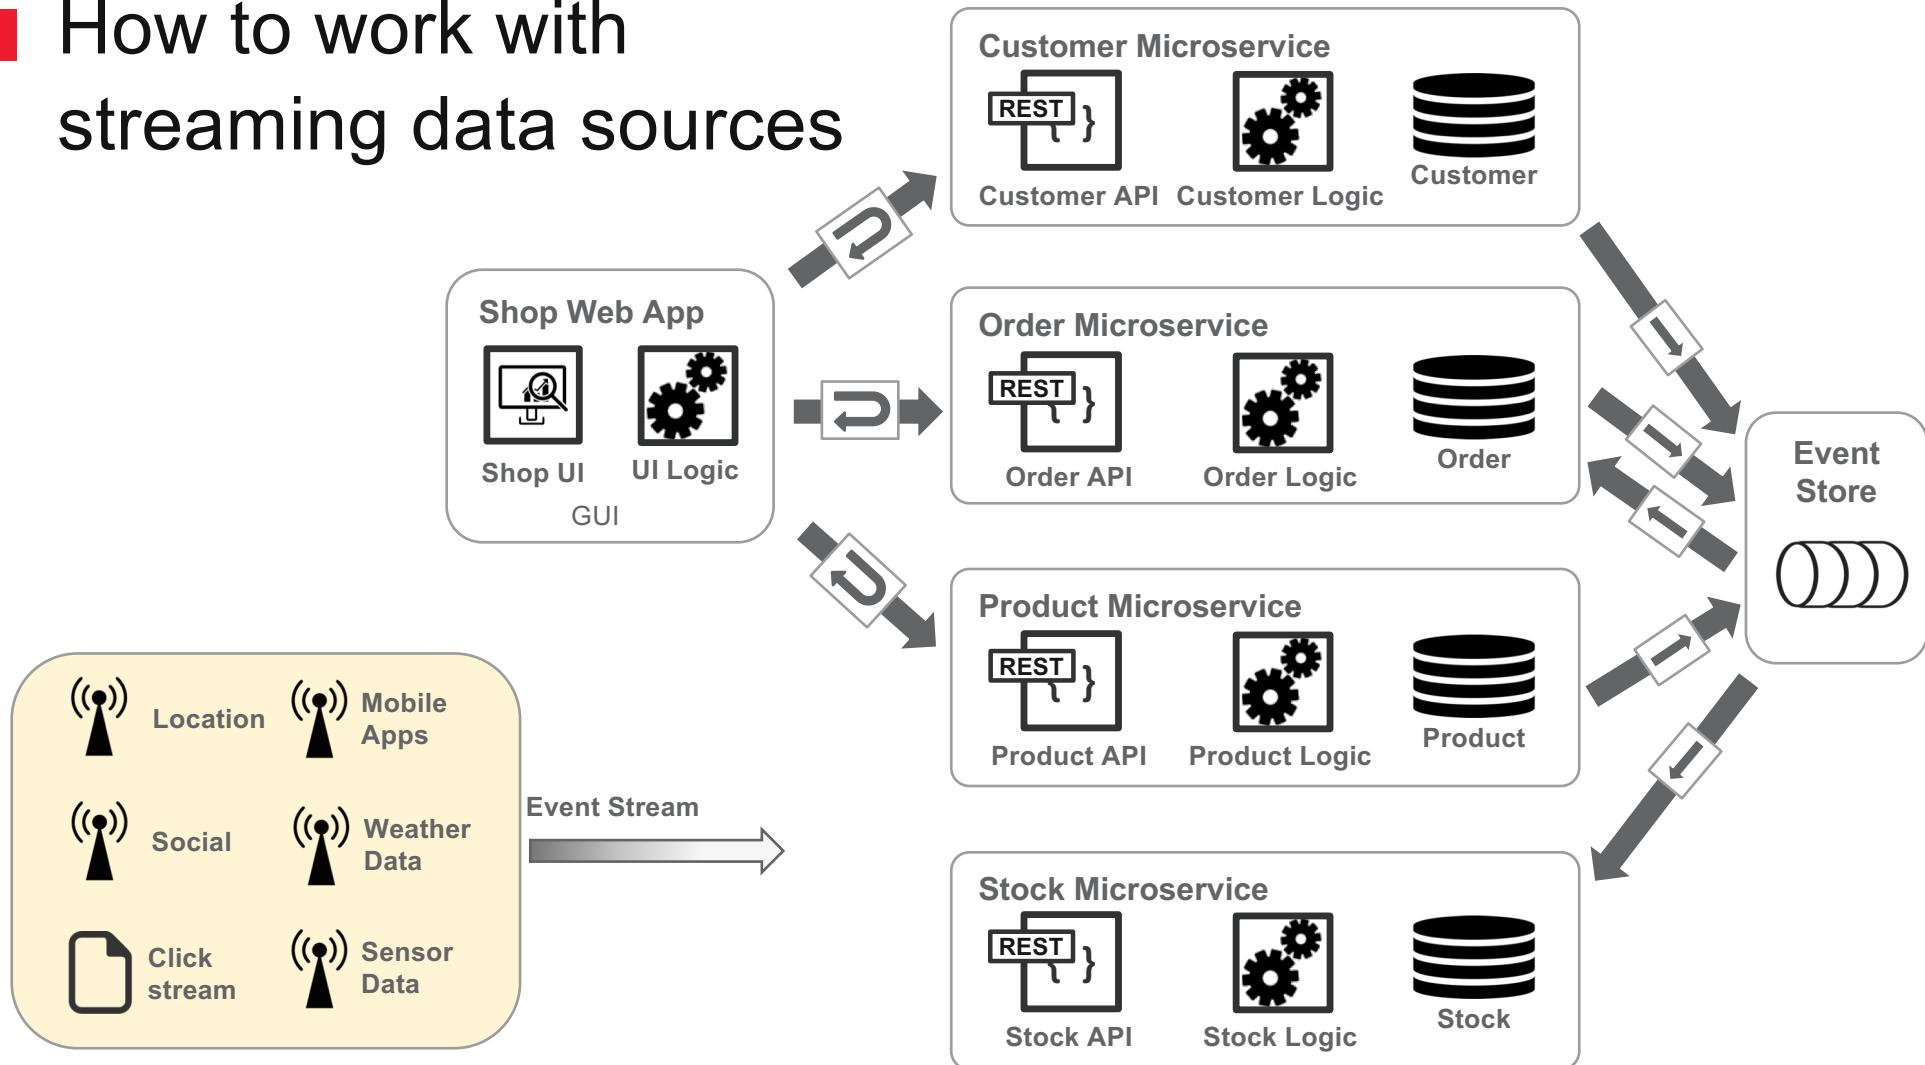
<https://github.com/gschmutz/various-demos/tree/master/event-driven-microservices>

■ Adding a new service by bootstrapping from Event Hub

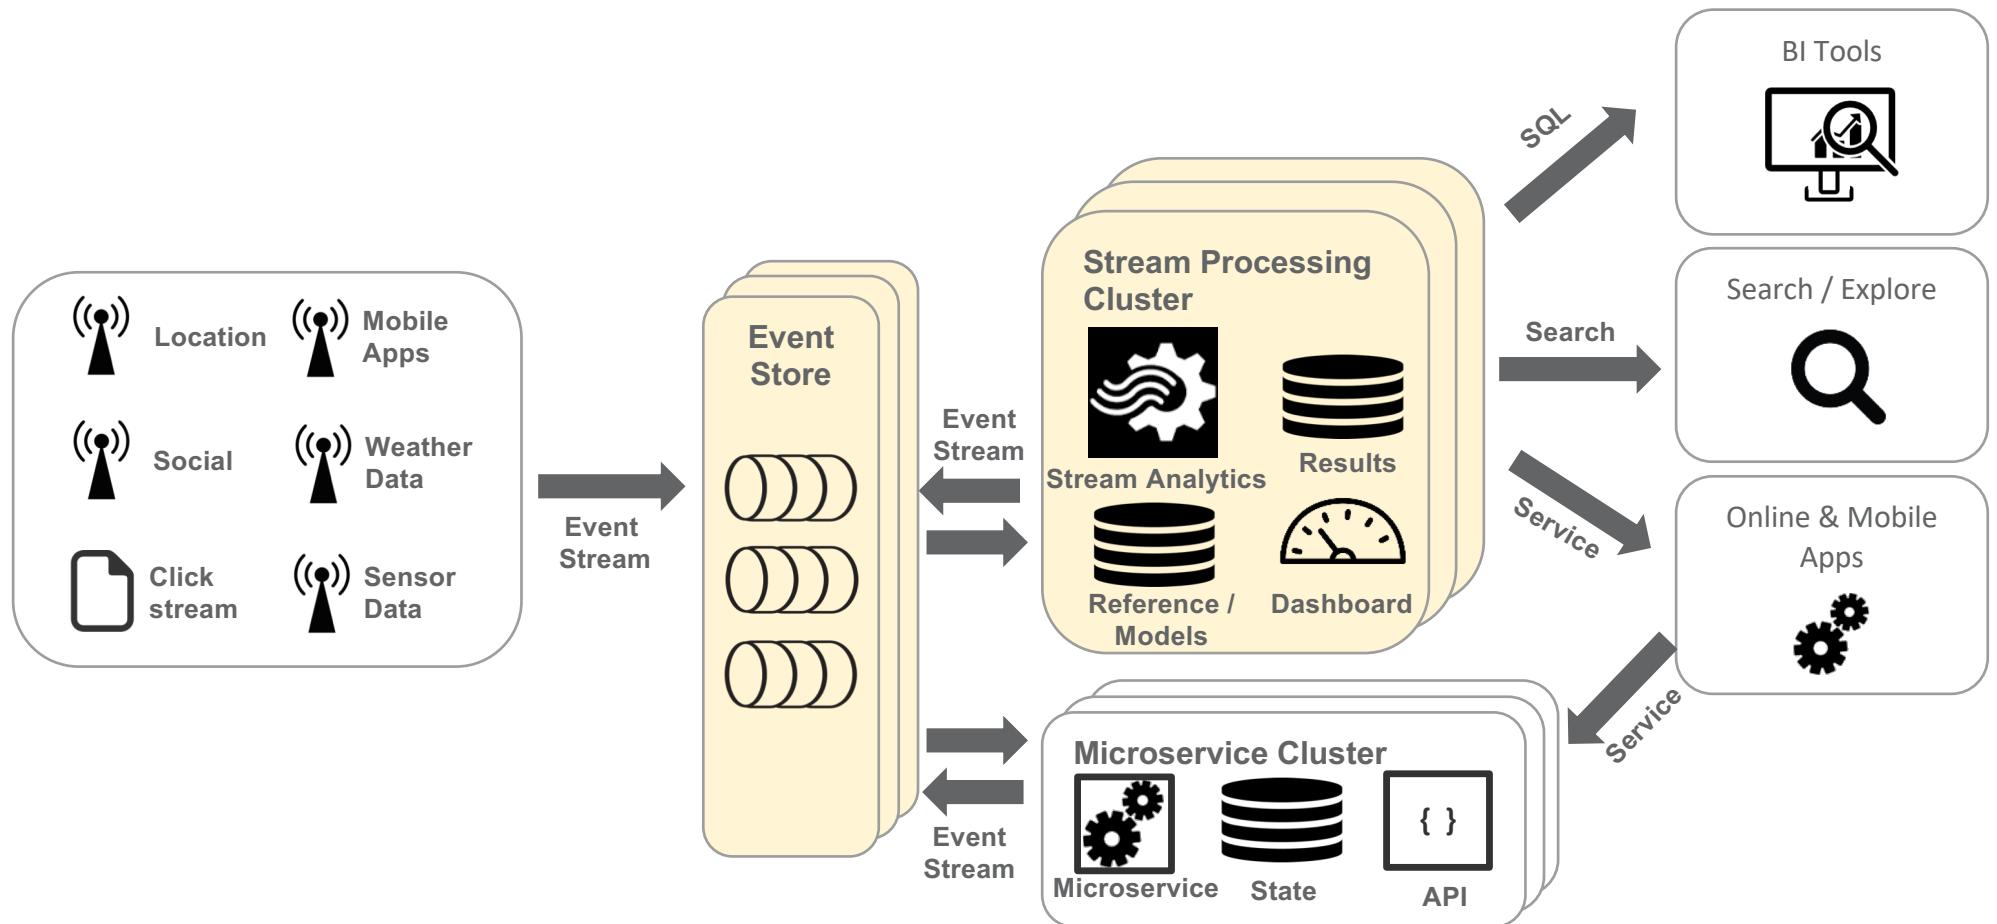


What about streaming sources?

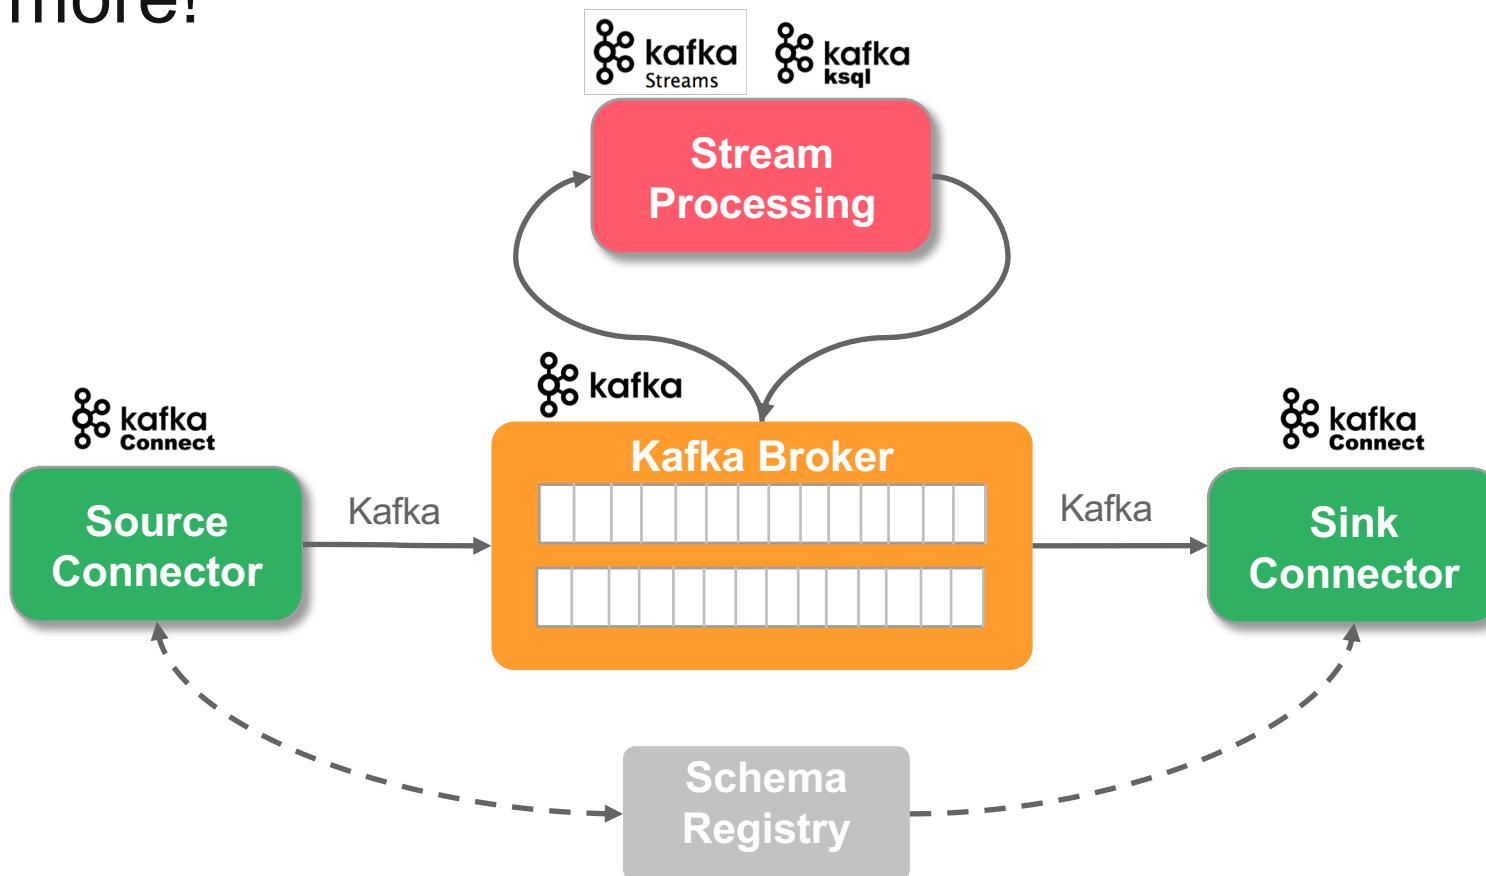
■ How to work with streaming data sources



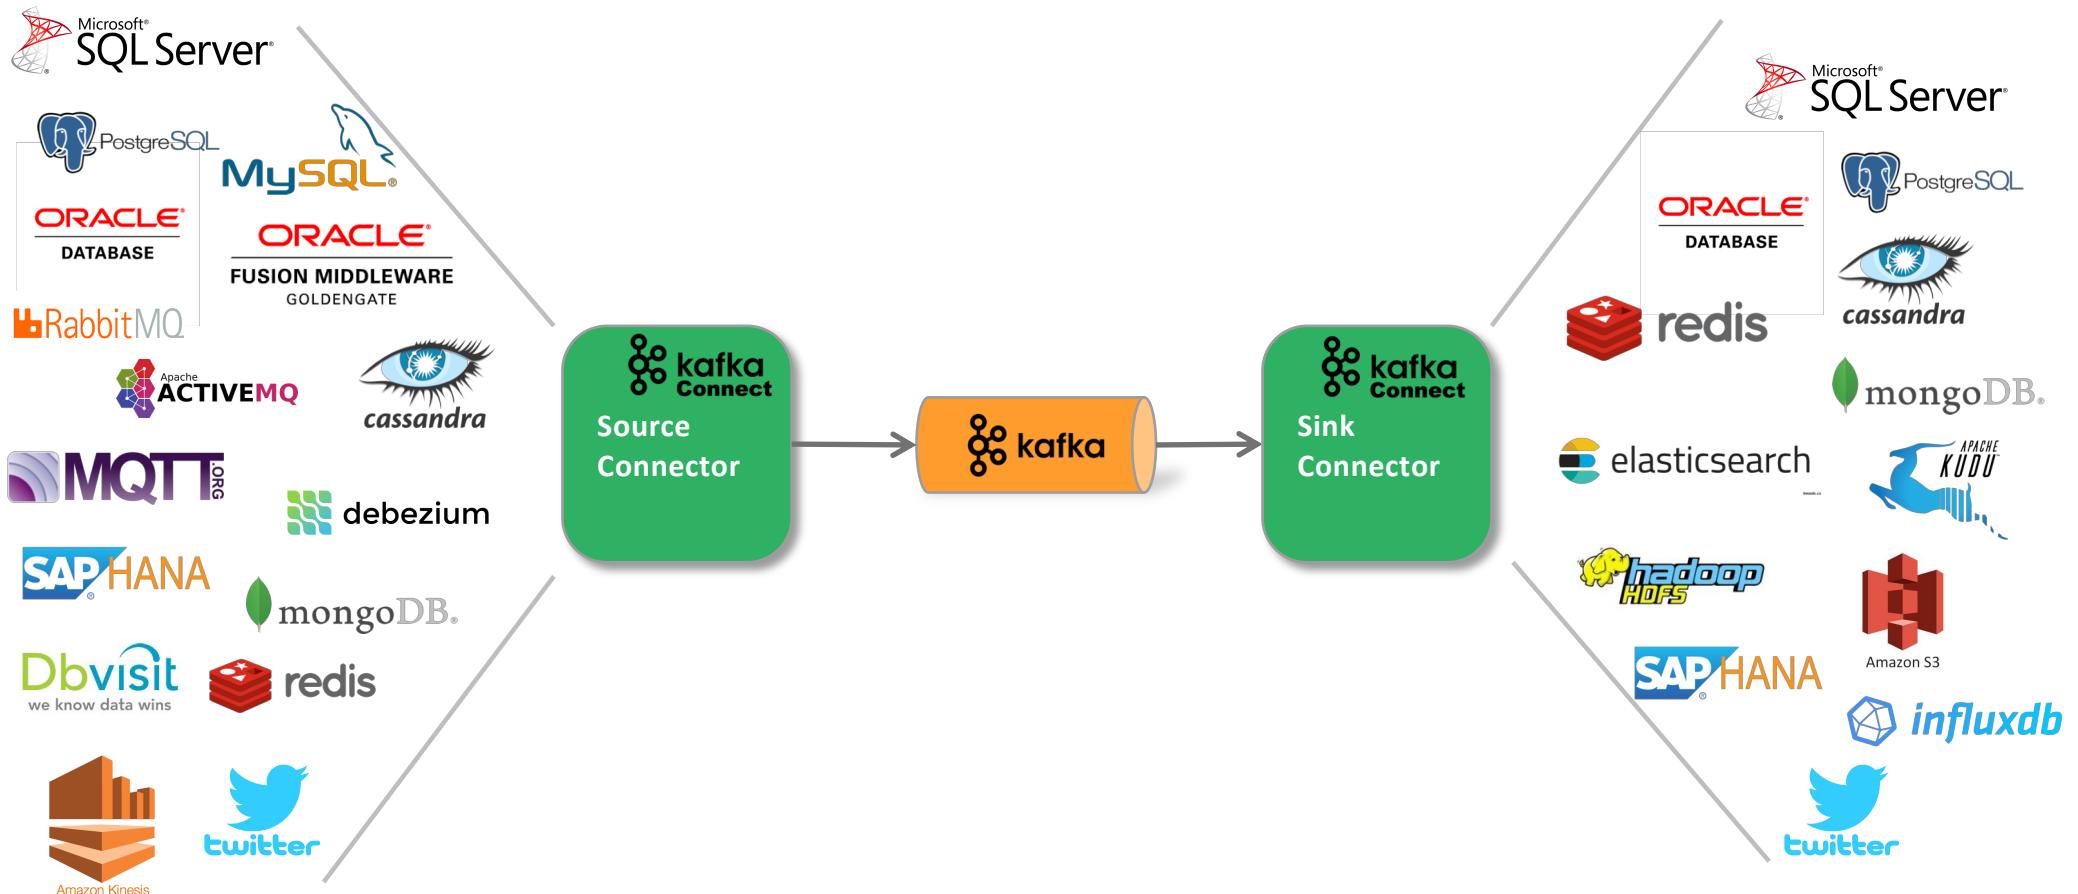
■ Streaming Processing & Microservices Architecture



Apache Kafka – scalable message processing and more!



Kafka Connect - Overview



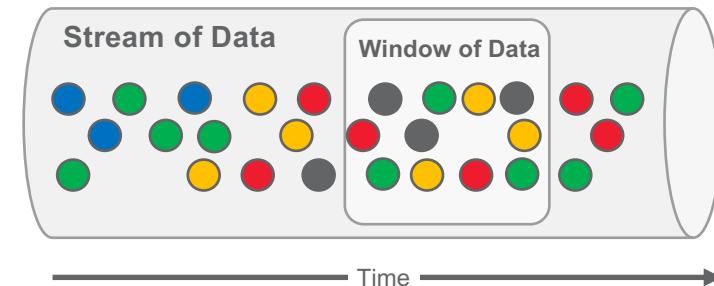
■ Stream Processing Capability - Windowing

Computations over events done using **windows of data**

Due to size and never-ending nature of it, it's not feasible to keep entire stream of data in memory

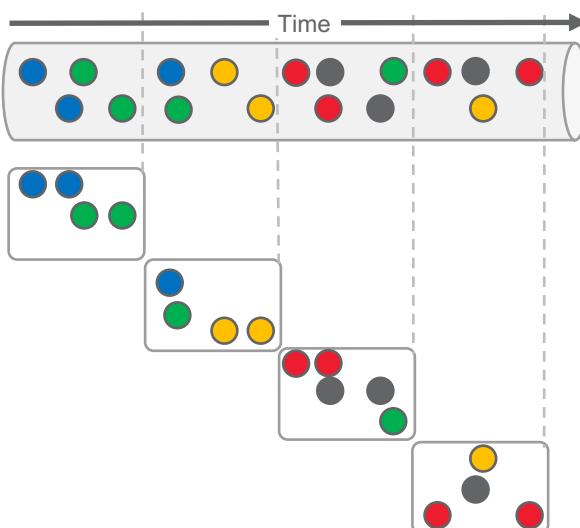
A window of data represents a certain amount of data where we can perform computations on

Windows give the power to keep a working memory and look back at recent data efficiently

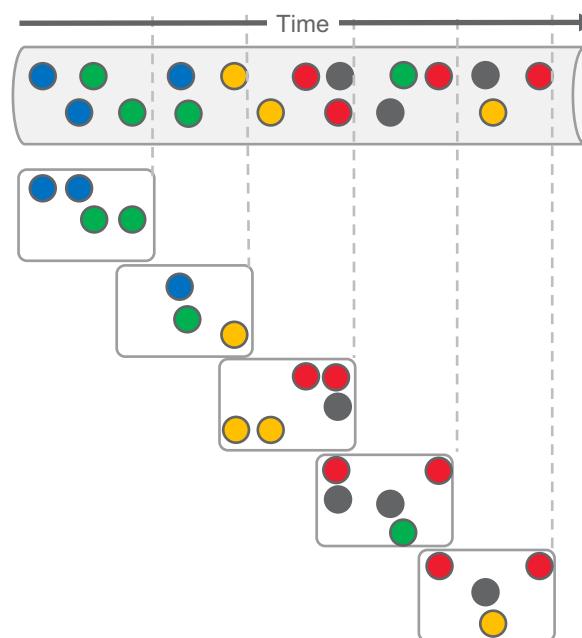


■ Stream Processing Capability - Windowing

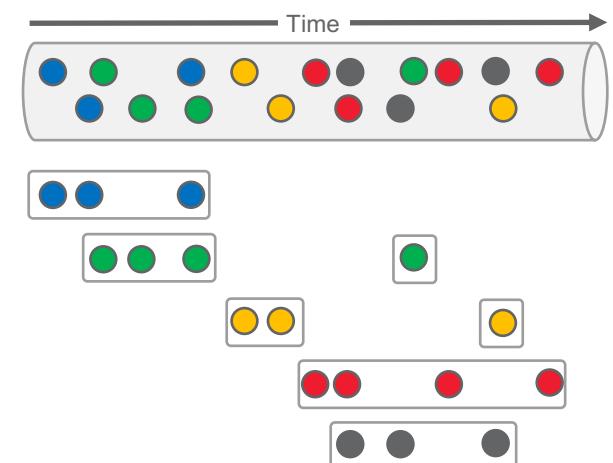
Fixed Window (aka Tumbling Window) - eviction policy always based on the window being full and trigger policy based on either the count of items in the window or time



Sliding Window (aka Hopping Window) - uses eviction and trigger policies that are based on time: *window length* and *sliding interval length*



Session Window – composed of sequences of temporarily related events terminated by a gap of inactivity greater than some timeout

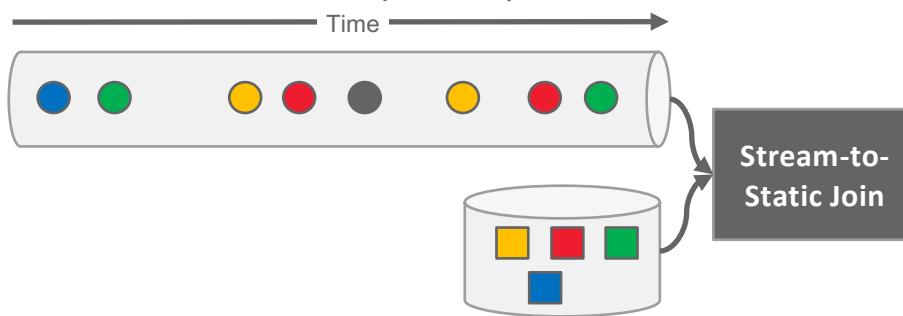


■ Stream Processing Capability - Joining

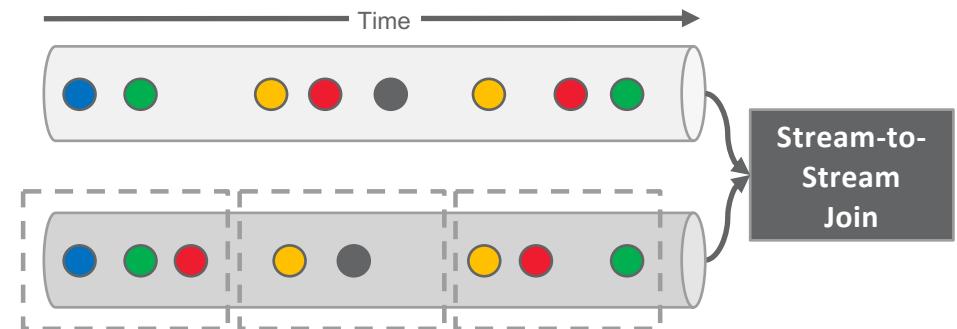
Challenges of joining streams

1. Data streams need to be aligned as they come because they have different timestamps
2. since streams are never-ending, the joins must be limited; otherwise join will never end
3. join needs to produce results continuously as there is no end to the data

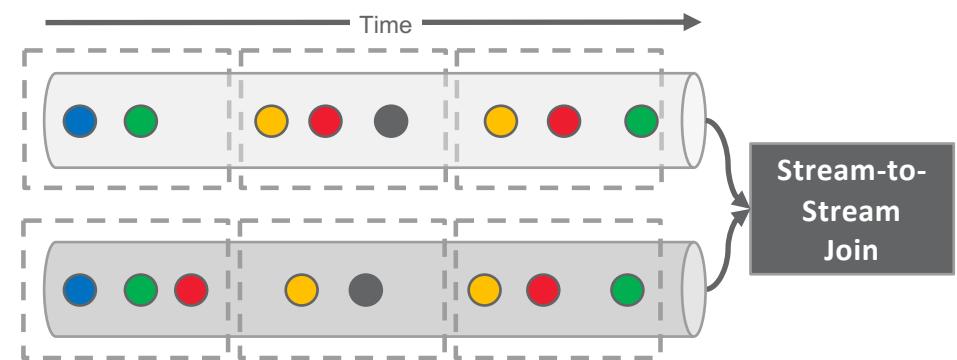
Stream to Static (Table) Join



Stream to Stream Join (one window join)



Stream to Stream Join (two window join)



■ Stream Processing Capability - State Management

Necessary if stream processing use case
is dependent on previously seen data or
external data

Windowing, Joining and Pattern

Detection use State Management behind the scenes

State Management services can be made available for custom state handling logic

State needs to be managed as close to the stream processor as possible

Options for State Management



How does it handle failures? If a machine crashes and the/some state is lost?

Kafka Streams - Overview

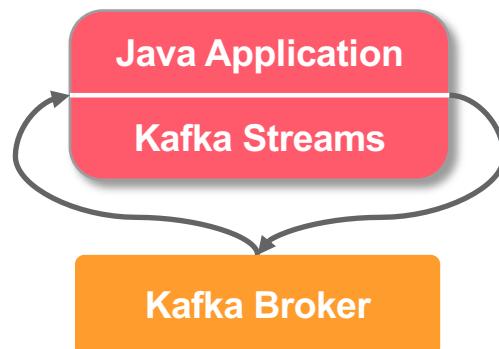


Designed as a **simple and lightweight library in Apache Kafka**

no other dependencies than Kafka

Supports **fault-tolerant local state**

Supports **Windowing** (Fixed, Sliding and Session) and **Stream-Stream / Stream-Table Joins**



Millisecond processing latency, no micro-batching

At-least-once and exactly-once processing guarantees

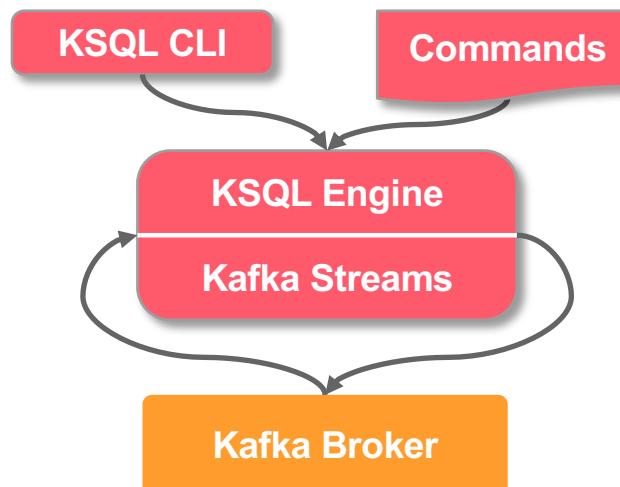
```
KTable<Integer, Customer> customers =  
    builder.stream("customer");  
  
KStream<Integer, Order> orders =  
    builder.stream("order");  
  
KStream<Integer, String> enriched =  
    orders.leftJoin(customers, ...);  
  
joined.to("orderEnriched");
```

■ KSQL - Overview

Stream Processing with zero coding
using SQL-like language

Built on top of Kafka Streams

Interactive (CLI) and headless (command
file)



STREAM and TABLE as first-class
citizens

- STREAM = data in motion
- TABLE = collected state of a stream

```

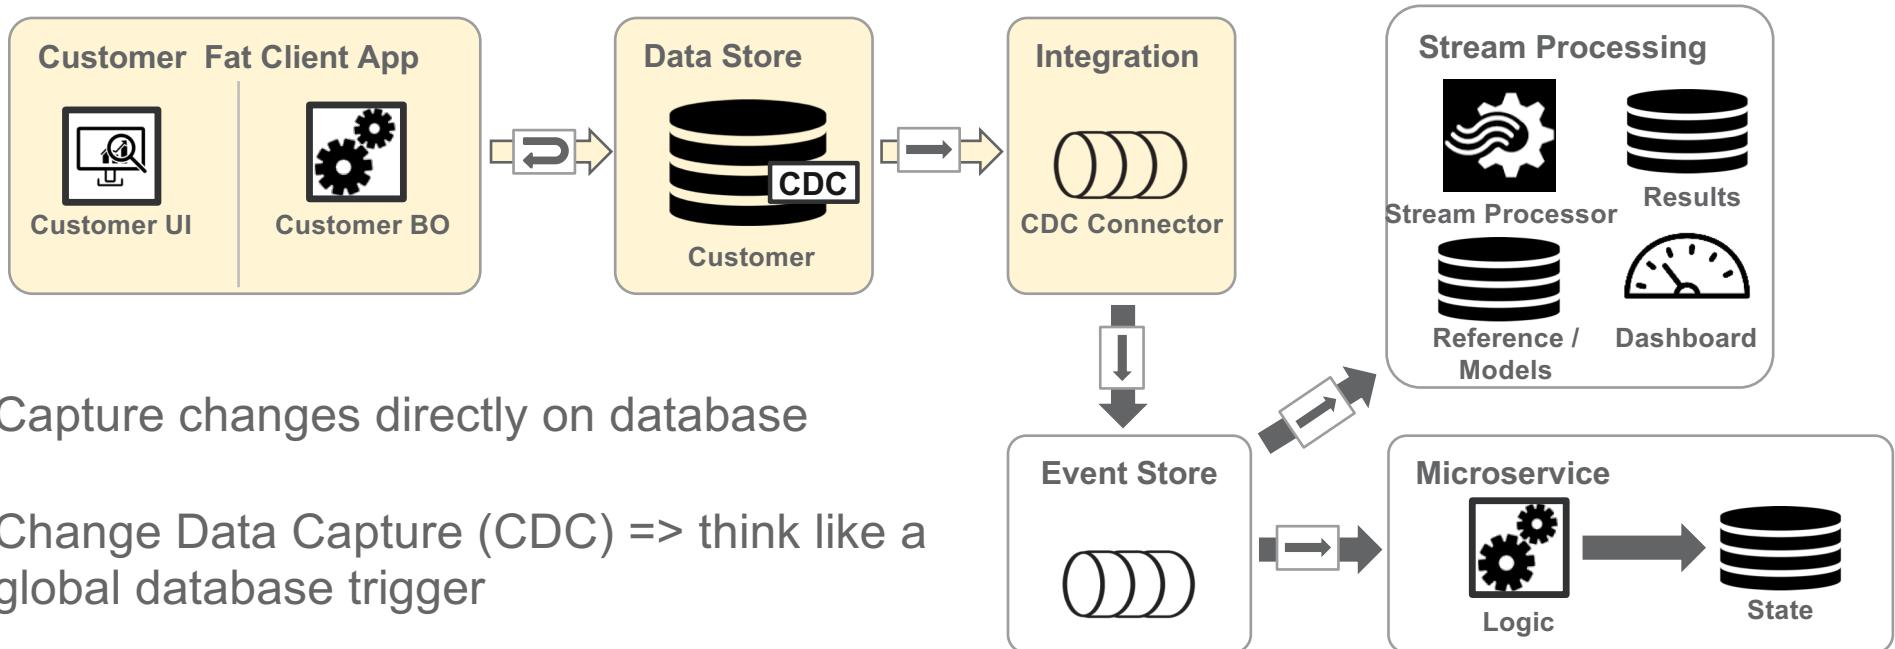
ksql> CREATE STREAM customer_s \
      WITH (kafka_topic='customer', \
            value_format='AVRO');

Message
-----
Stream created

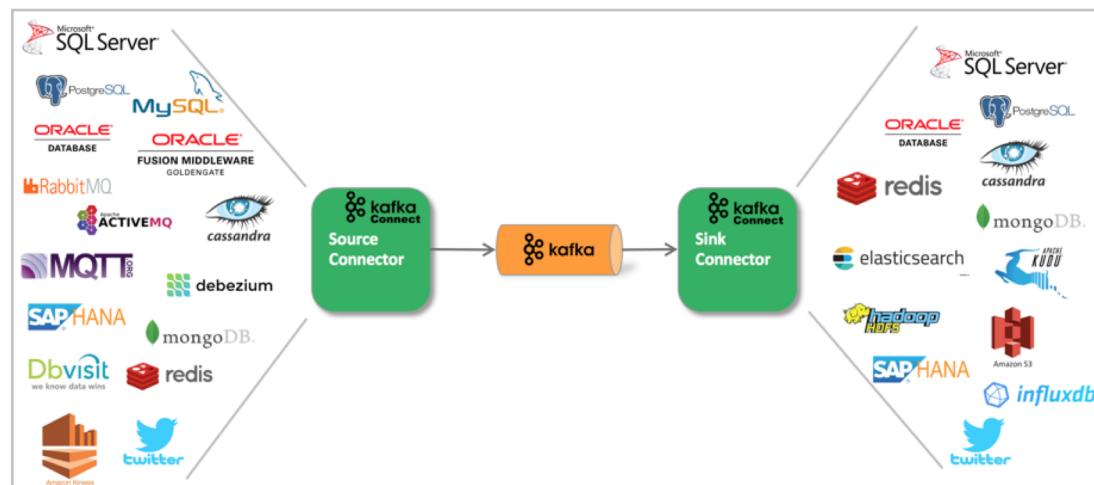
ksql> SELECT * FROM customer_s \
      WHERE address->country = 'Switzerland';
... 
```

What about integrating legacy applications?

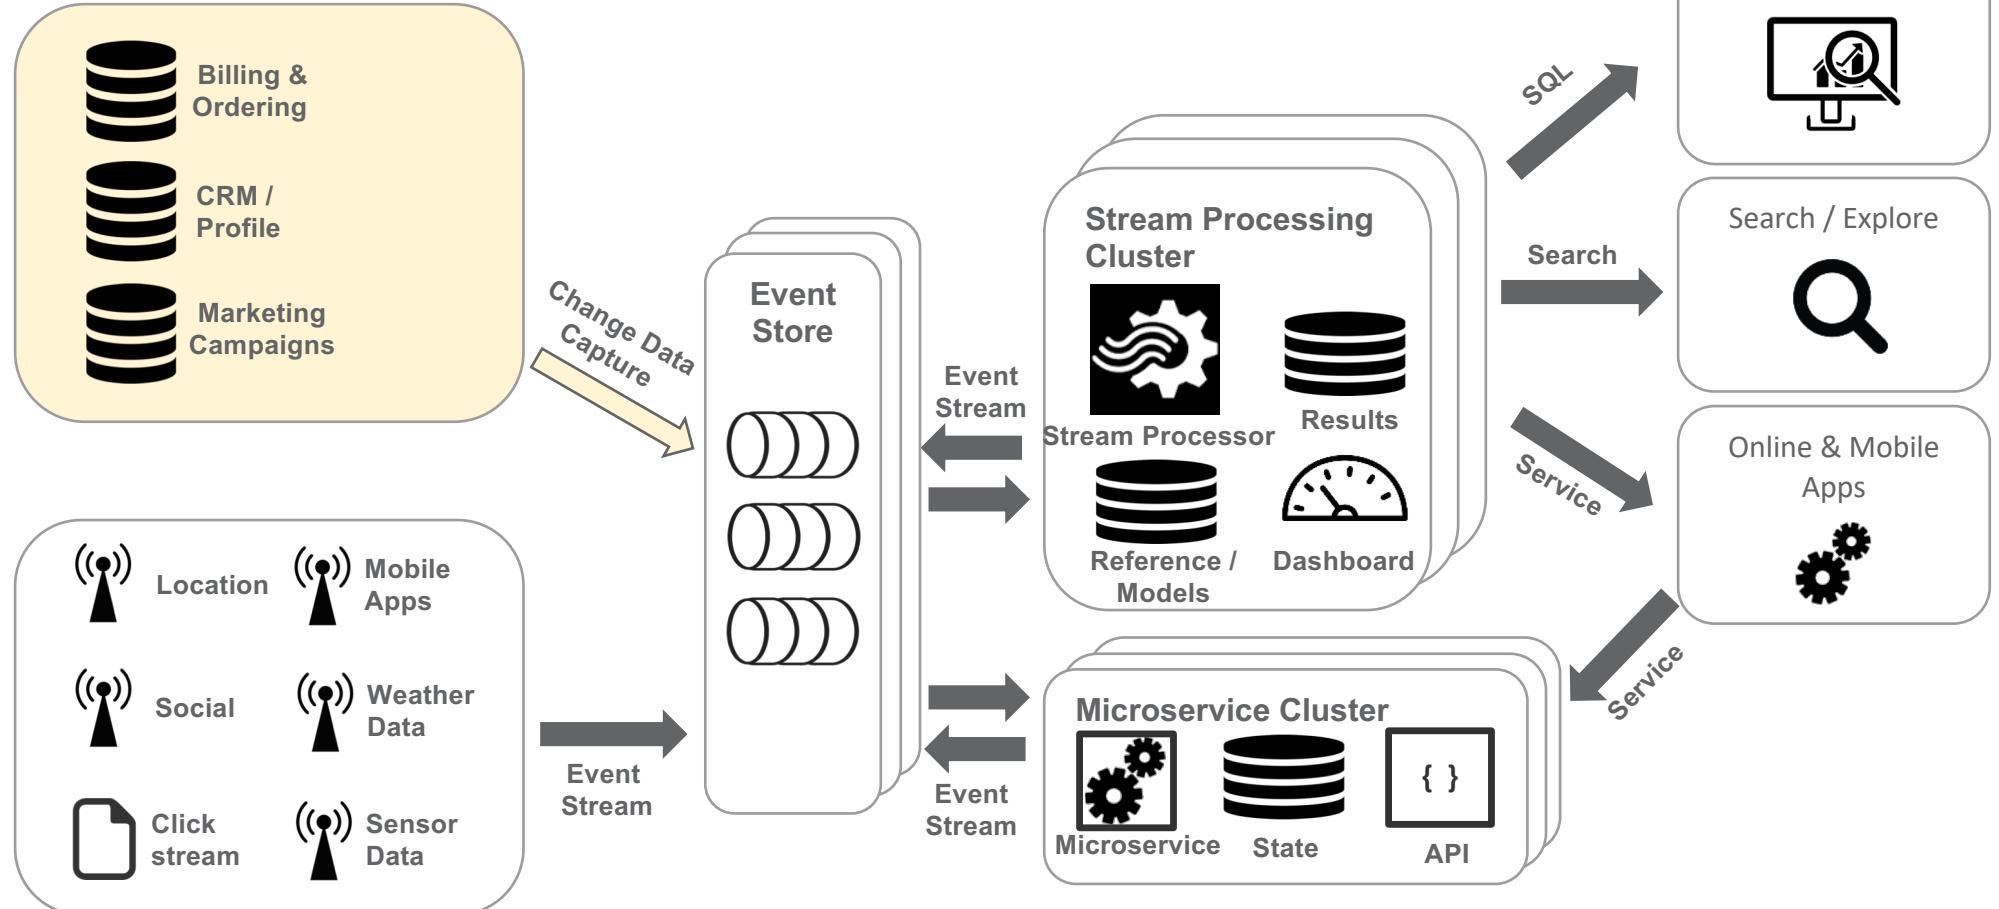
■ Integrate existing systems through CDC



■ Change Data Capture (CDC) with Kafka

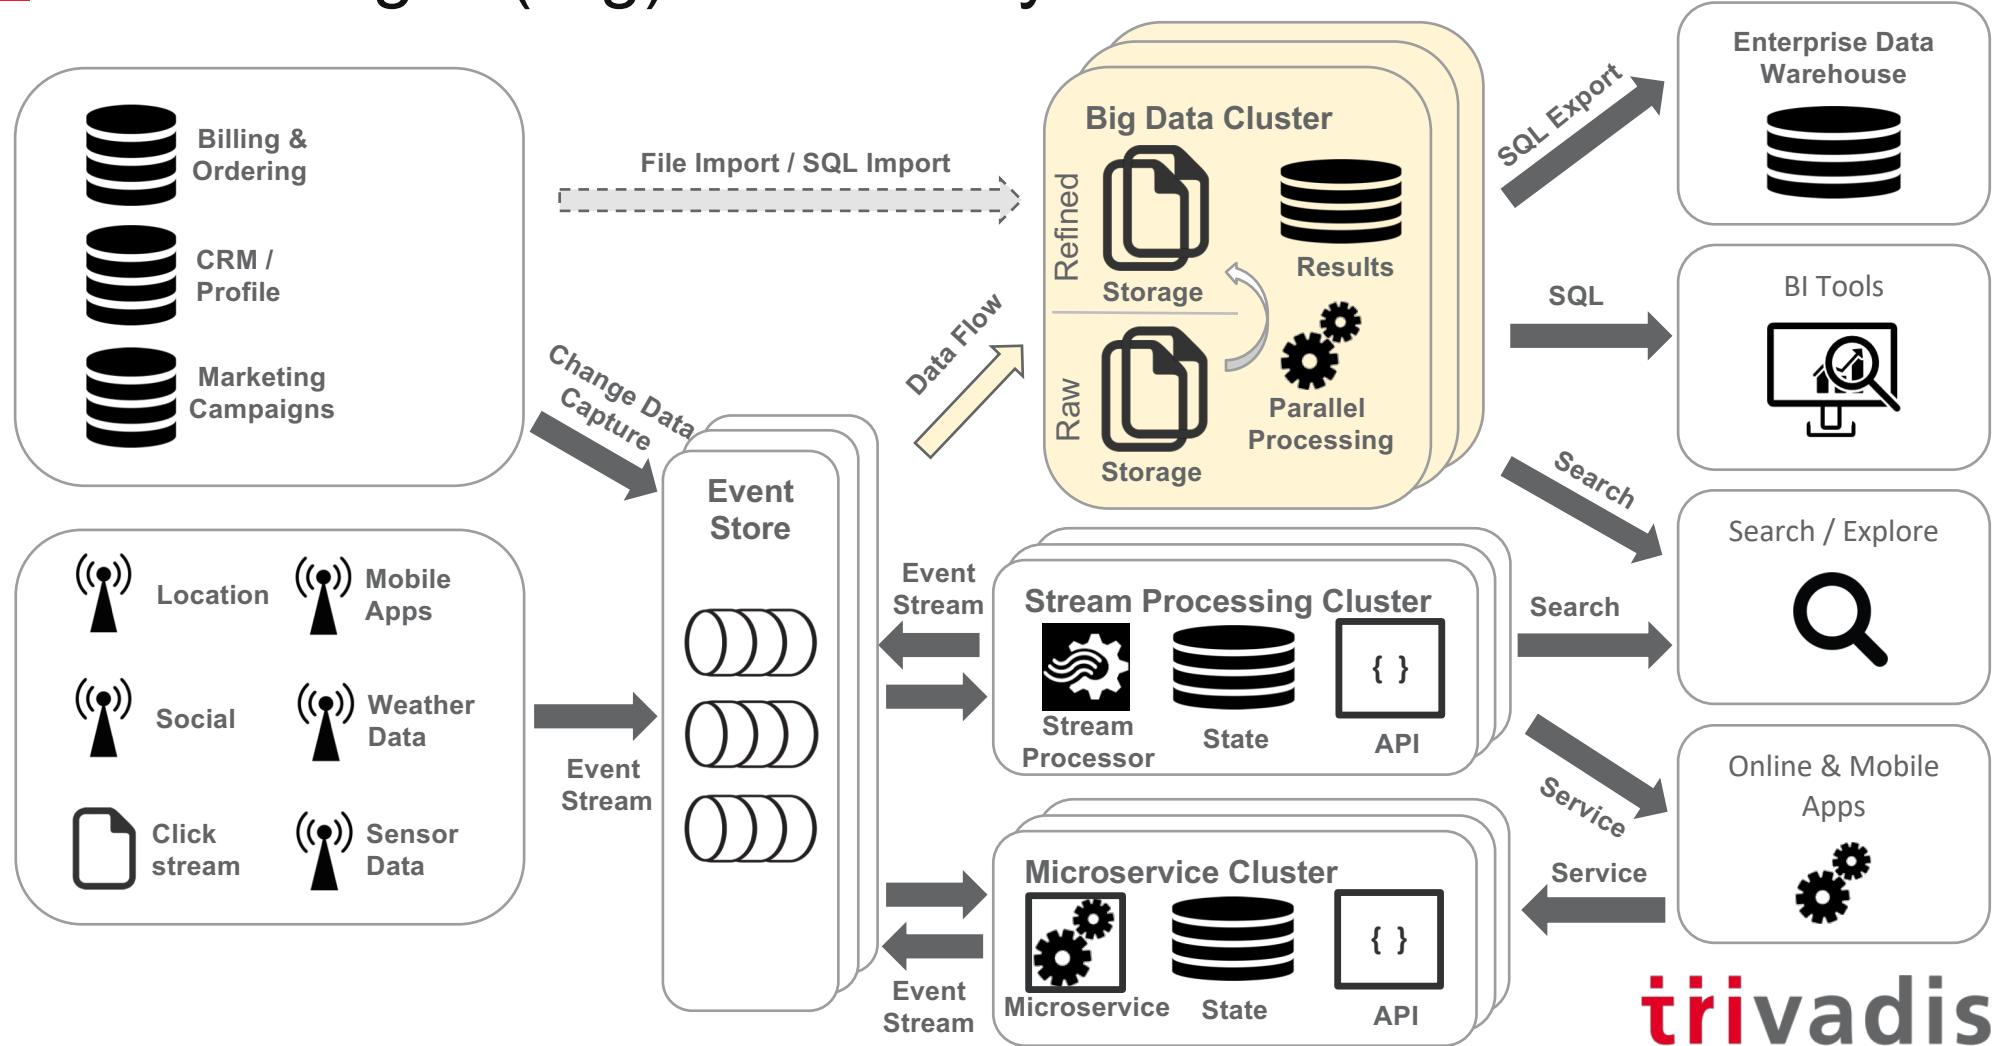


■ Integrate existing systems through CDC



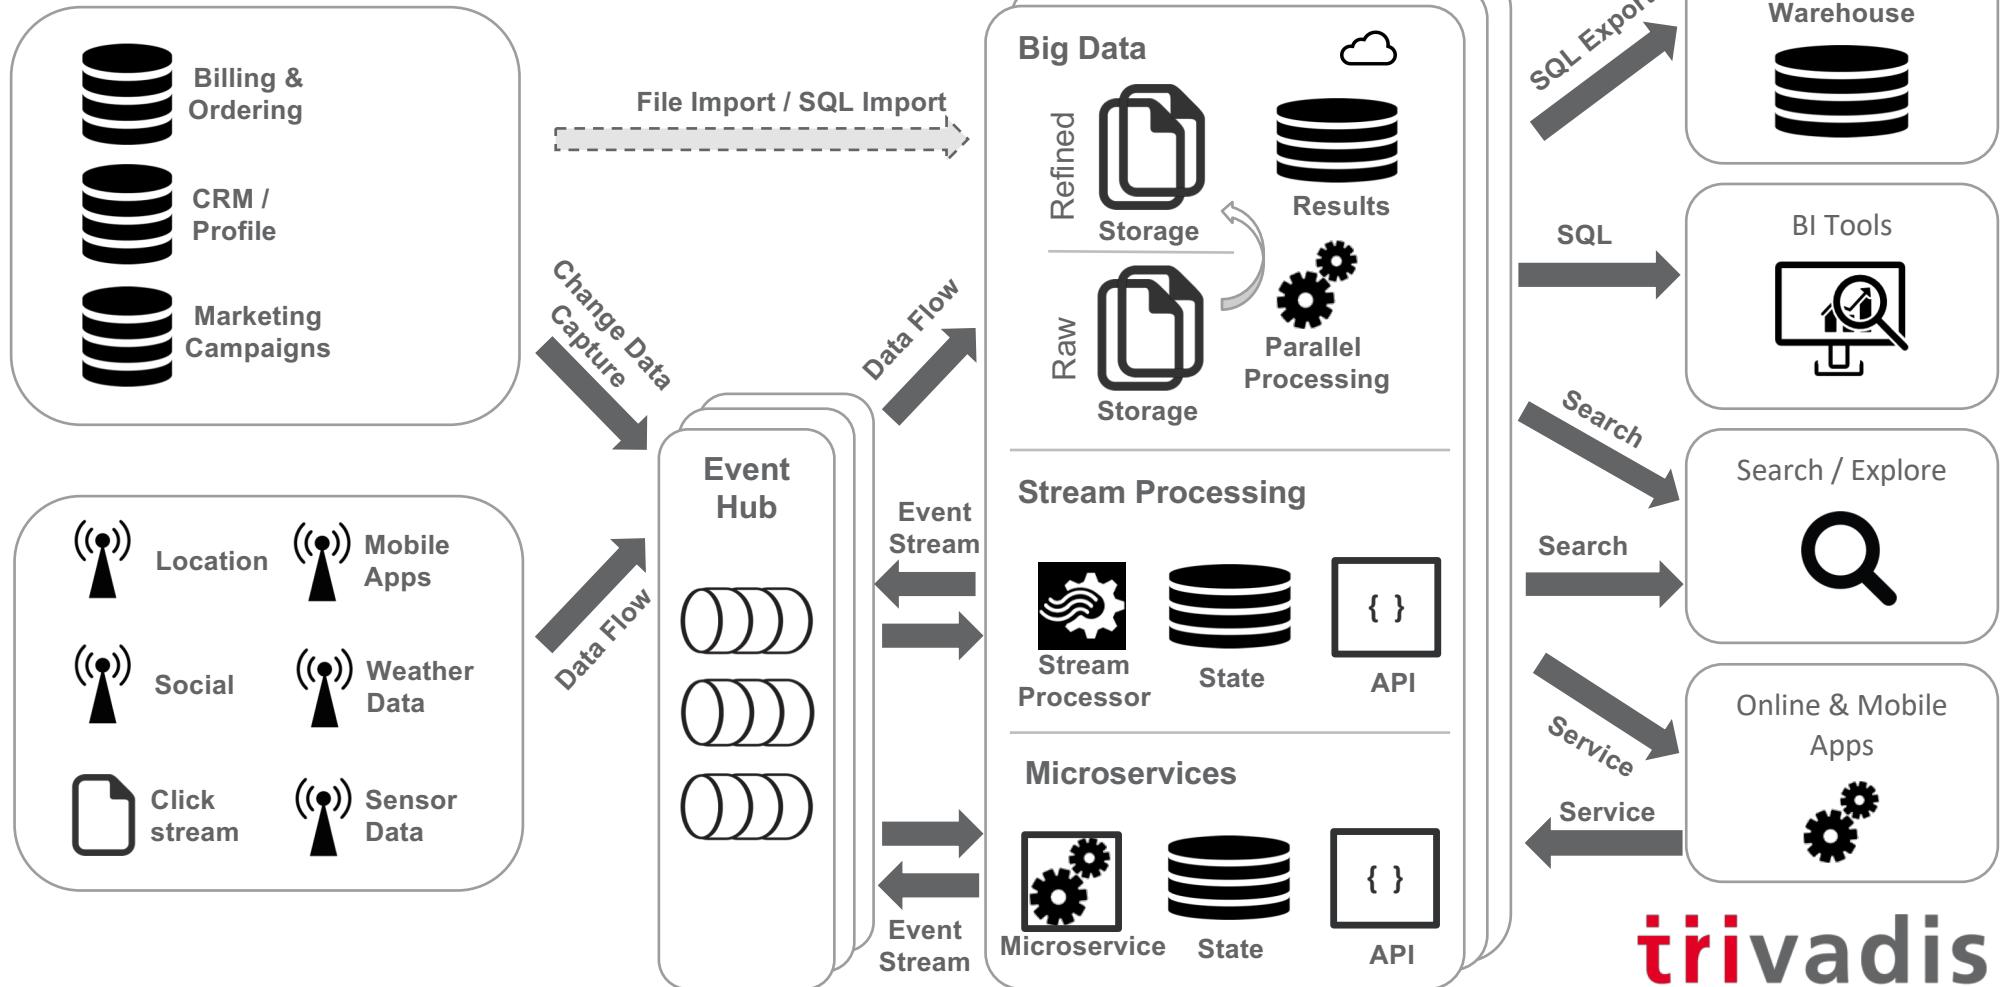
What about (historical) data analytics?

■ Streaming & (Big) Data Analytics Architecture



Summary

■ Summary



■ Summary

- service autonomy is key in a Microservices Architecture!
- not all communication need to be synchronous => separate into
 - commands
 - events
 - queries
- Kafka is well suited as an event broker / event store
 - brings many more interesting features beyond just “message passing”

■ References

Microservices Blog Series, Ben Stopford, Confluent:

- <https://www.confluent.io/blog/tag/microservices>

Apache Kafka for Microservices: A Confluent Online Talk Series:

- <https://www.confluent.io/landing-page/microservices-online-talk-series/>

Turning the database inside-out with Apache Samza, Martin Kleppmann, Con

- <https://www.confluent.io/blog/turning-the-database-inside-out-with-apache-samza/>

Event sourcing, CQRS, stream processing and Apache Kafka: What's the connection?, Neha Narkhede, Confluent:

- <https://www.confluent.io/blog/event-sourcing-cqrs-stream-processing-apache-kafka-whats-connection/>

Immutability Changes Everything, Pat Helland, Salesforce:

- http://cidrdb.org/cidr2015/Papers/CIDR15_Paper16.pdf

Commander: Better Distributed Applications through CQRS and Event Sourcing, Bobby Calderwood:

- <https://www.youtube.com/watch?v=B1-gS0oEtYc>

**Technology on its own won't help you.
You need to know how to use it properly.**



trivadis
makes **IT** easier. 